

Grounded Harness Brief

Executive Brief — Grounded Harness Constraint Transport (2026-07-20)

We evaluated Constraint Transport and Grounded Statecharts in experiments/grounded_statecharts under a live, matched-budget harness on gpt-4.1-mini. Early labeled-prompt D2 looked strong (CT joint_success $\delta = +1.0$; GS false_completion $\delta = -0.167$), but a weak-prompt ablation that removed condition labels collapsed both effects to 0.0, proving those gains were prompt/scaffold artifacts. After redesigning to name-free default prompts, harness enforcement in condition_policy.py (G3 artifact repair; external envelope capability narrowing), and evidence-based scoring, CT joint-success returned to $\delta = +1.0$ on the harness-v2 name-free ablation, held-out D2 (144 episodes; CI [1.0, 1.0]), and D3 confirmatory (120 episodes, 5 nested repeats; CI [1.0, 1.0]), while GS stayed narrowed. The useful claim is operational: external guards recover joint success after capability widenings under name-free prompts via harness enforcement—not that the model learned constraints. That matters for recursive agent stacks because it relocates reliability from label-sensitive prompting to enforceable child contracts, with public sanitized datasets at results/d2_pilot_public/ and results/d3_ct_confirmatory_public/. Next: OOD wording/depth probes for CT, broader CHS seals beyond orchestration, and independent HU corpora before any unlearning claim.