


The Lineage and the Trajectory

*A History of the Ideas Behind the Research Program —
Where They Came From, and How the Program Grew*



A companion to *The Mathematics of Constraint*.

Part I traces the decades of science the program stands on.

Part II reconstructs how the program itself evolved, paper by paper.

Why History Matters Here

No idea is born from nothing, and this program less than most. Its central conjecture — that adaptive systems keep rediscovering geometric language because geometry is the portable language of constraints — is itself a claim *about history*: that a dozen independent fields, working for decades in ignorance of one another, kept arriving at the same shapes. You cannot judge that claim without knowing the fields. And you cannot understand any individual result without knowing which tradition it extends, tests, or departs from.

This book supplies both histories. **Part I** is the *external* lineage: the intellectual ancestry the program draws on — cybernetics and autopoiesis, information theory and rate-distortion coding, group theory and topology, causal inference and reinforcement learning — with the key figures, the seminal ideas in plain language, and exactly how the repository uses each. **Part II** is the *internal* trajectory: how the program itself developed, from a single founding question through a numbered chain of some thirty-five experiments, its "correction chain" of self-discovered mistakes, its pivots and abandoned directions, and its current frontier.

How to read this book

You do not need the mathematics to follow the history; where a technical idea appears, it is explained in a sentence, with a pointer to the math primer for depth. Read Part I to understand what the program inherited and Part II to understand what it did with the inheritance. A recurring **Verdict** box gives a mature assessment — what is genuinely novel versus what is skilled recombination — because a history that only celebrates is not a history.

Contents

Part I · The Lineage

- 1 · The Convergence Puzzle
- 2 · Cybernetics, Autopoiesis, and the Systems Tradition
- 3 · Information, Compression, and Constraint
- 4 · Symmetry, Topology, and the Brain
- 5 · Causality, Control, and Agency

Part II · The Trajectory

- 6 · The Founding Question and the Three Threads
- 7 · The Numbered Spine and the Correction Chain
- 8 · Arc 2, the Benchmark, and the Phases
- 9 · The Reframe and the Frontier
- 10 · The Meta-Story: Machine-Made Science

PART I

The Lineage

Before the program's own story comes the story it inherited. This part traces the traditions its ideas descend from — not as a dry bibliography, but as a living intellectual family tree, so that when you meet "weakness" or "the metric deforms under concern" you hear the centuries of thought standing behind the phrase. We begin with the puzzle that motivates the whole enterprise.

CHAPTER 1

The Convergence Puzzle

1.1 The observation that started it

Here is the pattern the program exists to explain. Over the twentieth and twenty-first centuries, a striking number of independent fields — mathematics, neuroscience, linguistics, control theory, machine learning, theoretical biology — kept reaching, apparently without coordination, for the same kind of language to describe meaning, thought, and agency: *geometric* language. Distances and neighborhoods. Attractors and basins. Manifolds, boundaries, directions, curvature, regime transitions. A neuroscientist describing how the brain codes space, a linguist describing how word meanings relate, an AI researcher describing a neural network's internal states, and a mathematician describing a dynamical system would, more and more, find themselves drawing the same pictures.

The program's founding document states the puzzle as its central question:

The manifesto, verbatim

"Why do independently developed systems of thought and computation keep converging on geometric descriptions of meaning, agency, and intelligence, and can we predict when that geometry is merely a passive representation versus when it becomes an active, self-maintaining attractor regime?"

And the compressed thesis: "adaptive systems keep rediscovering geometric language because geometry is the portable language of constraints."

1.2 The proposed answer: a shared problem, not a shared substrate

Why would unrelated fields converge? The program's answer is elegant and worth stating carefully, because it is the hypothesis every later chapter tests: they converge *not because they study the same stuff*, but because they all solve *the same problem*. That problem is compression under constraint. Any finite system — a brain, a network, a language, an organism — cannot keep every possible state of the world separate; it has limited capacity. So it must fold the world down into a lower-dimensional organization that keeps the *useful* distinctions and throws the rest away. And geometry, the program argues, is simply the natural language for that folding: it is the mathematics of "which things are near, which are far, which are reachable, which are stable" — exactly the relations a compressed-but-useful representation must preserve.

Intuition pump — why every subway map looks the same

Subway maps from Tokyo, London, and New York look eerily alike: colored lines, dots for stations, distances distorted, geography sacrificed. No cartographer copied another; they converged because they all solved the same problem — *show a rider which stops connect to which, and throw away everything else*. The true geography (the constraint of the physical city) is folded down to the one thing that matters for the task (connectivity), and the result is a shared visual language. The program's claim is that minds, networks, and organisms are all drawing subway maps of reality, and that is why they rhyme. The geometry is not proof they are the same thing; it is evidence they face the same compression problem.

1.3 The disciplined form of the claim

What keeps this from being mysticism is a single sentence of restraint that recurs throughout the program: "*Geometry is not proof of mind. But recurring geometry is evidence of a shared constraint problem.*" This is the crucial move. It would be easy — and wrong — to leap from "brains and neural networks both use geometric codes" to "neural networks have minds." The program explicitly refuses that leap. The convergence is evidence of a common *problem structure*, not a common inner life. The sharpened research question is therefore not "are these systems the same?" but "why does the same *form of explanation* keep winning?" — a question that can actually be investigated.

1.4 The map of what feeds in

To explain a convergence, you must know what is converging. The rest of Part I lays out the tributaries — the traditions whose independent arrival at geometric-constraint language is the phenomenon to be explained, and whose tools the program borrows to study it:

Tradition	What it contributes	Chapter
Cybernetics & autopoiesis	self-maintenance, feedback, viability, the living boundary	2
Information theory & MDL	compression, description length, the simplicity principle weakness challenges	3
Rate-distortion & efficient coding	how finite capacity should allocate resolution — the "reward deforms the metric" law	3
Group theory & equivariant ML	symmetry as the backbone of weakness and generalization	4
Topology & neural manifolds	the torus code, shape as a substrate-independent invariant	4
Causal inference	interventions, identifiability, the do-operator, the self/world gauge	5
Reinforcement & homeostatic learning	agents as viability-maintainers, value, planning	5

Verdict — is the convergence thesis itself novel?

A mature reader should note up front: the grand convergence thesis is a *synthesis of others' ideas*, and the program says so. That many fields reach for geometry has been observed before — by the efficient-coding neuroscientists, by the "Platonic Representation Hypothesis" in machine learning (Chapter 4), by the enactivists. The program's original contribution is not the observation but the attempt to make it *predictive*: to name the specific constraint (finite capacity), the specific mechanism (weakness / symmetry-compatible compression), and the specific threshold (passive representation becoming active agency), and then to test these in experiments. Whether that sharpening succeeds is the subject of Part II. The thesis as inspiration is borrowed; the thesis as a falsifiable research program is the program's own.

1.5 What to carry forward

Many independent fields converged on geometric language for mind and agency; the program explains this as a shared compression-under-constraint problem, not a shared substrate; it disciplines the claim with "geometry is evidence of a shared problem, not proof of mind"; and its originality lies in trying to make an old observation predictive rather than in the observation itself.

CHAPTER 2

Cybernetics, Autopoiesis, and the Systems Tradition

The program's ideas about *agency* — self-maintenance, viability, repair, the living boundary — descend from a lineage that began in the 1940s and has never quite entered the mainstream. This chapter traces it, because more than any other tradition it supplies the program's answer to "when does a representation become an agent?"

2.1 Ashby and the birth of ultrastability (1948–1960)

W. Ross Ashby, a British psychiatrist, asked a deceptively simple question: how does an organism stay alive in a world that keeps trying to knock it out of its viable range? His answer, in *Design for a Brain* (1952/1960), introduced two ideas the program uses directly. **Requisite variety**: to control a disturbance, a regulator must have at least as much variety in its responses as the disturbance has in its assaults — you cannot stabilize what you cannot match. And **ultrastability**: when a system is pushed outside its viable zone, it does not merely resist; it *reorganizes its own internal parameters* until stability returns. The thermostat that can only heat and cool is stable; the creature that can rewire itself when heating and cooling stop working is *ultrastable*.

In the repo — "Ashby ultrastability in the small"

The autopoietic-control experiments make Ashby measurable. A trained model is deliberately damaged with random weight noise (accuracy drops to 0.45), then allowed a handful of self-supervised update steps. It climbs back to 0.965 — recovering its competence by reorganizing itself from its own input stream. The paper names this exactly: "Ashby ultrastability in the small," and cites *Design for a Brain* as "the conceptual ancestor of the repair experiment." Sixty years of theory becomes a measured recovery curve. The amount of slack a system has in its lower layers, the program finds, predicts how well it repairs — which is the program's own quantitative extension of Ashby's qualitative idea.

2.2 Wiener, and the word "cybernetics"

Norbert Wiener gave the field its name in 1948, from the Greek for "steersman." Cybernetics is the science of *feedback* — of systems that sense their own error and act to reduce it. The steam-engine governor that pinches the throttle as the engine speeds up is the ancestral image: no intelligence, just a loop converting "too fast" into "less fuel." The program's homeostatic agents are Wienerian feedback loops at heart — an internal energy variable, a decay that pushes it down, actions that push it back up — though, tellingly, the repository invokes Ashby by name far more than Wiener. It is a lineage the program uses without always crediting, which a careful historian should note.

2.3 Maturana, Varela, and autopoiesis (1970s–1980s)

The Chilean biologists Humberto Maturana and Francisco Varela pushed the systems tradition to its most radical claim: a living system is defined not by its *parts* but by a **self-producing organization** — an *autopoiesis* (self-making) — that continuously rebuilds its own boundary and components. A cell is not a bag of chemicals; it is the ongoing process that makes the bag that contains the chemicals that make the bag. From this grew **enactivism** (Varela, Thompson, Di Paolo): the view that cognition is not passive representation but active *sense-making* by a self-

maintaining organism, and that meaning is always organism-relative — the same world matters differently to different living things.

Verdict — a lineage claimed more than earned

Here a mature reader must be precise, because the program is. Its experiments invoke autopoiesis, but the strict Maturana–Varela criterion — a system that *produces its own components and boundary* — is *not met*. The program's "repair" is a model updating its weights via labeled gradient descent, not a system regenerating itself from within. The papers admit this directly: "the system does not yet generate new policy, it only updates head weights." What the program has is *homeostasis* (staying in range) and *ultrastability* (reorganizing to stay in range), which are real and measured — but not full autopoiesis. Notably, Maturana and Varela are cited only in the synthesis papers, never in the primary experiments, which route autopoiesis entirely through Bennett's more recent formalization. The tradition supplies the *vocabulary and the aspiration*; the experiments deliver something more modest and should be read that way.

2.4 Bennett: the modern hub (2023–2026)

If one contemporary figure is the program's true intellectual center, it is Michael Timothy Bennett, whose doctoral work at ANU the repository leans on more than any other single source. Bennett recombined the older systems tradition with information theory into a set of sharp, formal claims the program then tests:

- **Weakness over simplicity.** The best hypothesis is not the shortest but the *weakest* — the one compatible with the most possibilities while still fitting the data (Chapter 3 and the math primer develop this). This is the program's flagship idea, inherited whole from Bennett and turned into an experiment.
- **The Autopoietic Theorem** (Bennett & Suzuki): life-like organization is *derivable* from weak axioms — that things change, that information capacity is finite, that low-level conditions are stable. The program operationalizes its two testable consequences: a *viability buffer* (how many futures a policy keeps open) and the *Law of the Stack* (a layer's flexibility is capped by the slack beneath it).
- **The tapestry of valence** and **first-order selves**: that a system represents not one drive but a pattern of effects on many internal variables, and that selfhood comes in layers. The program's vector-valued concern experiments test this directly.

To a large degree, then, the program is an **empirical test-bed for Bennett's theory** — taking claims that Bennett argues philosophically and asking whether they hold up as measured results in small neural systems. This is an honorable and unusual role, and it is worth seeing clearly: many of the program's most distinctive terms are not its own coinages but Bennett's, put to the test.

2.5 The virtual governor and Levin (2021–2026)

A newer strand comes from Michael Levin's work on *basal cognition* and multiscale agency — the idea, from developmental biology, that goal-directed, problem-solving behavior appears at every scale of life, from cells to tissues to organisms, and that a "self" is a provisional, nested, scale-relative construction. The program borrows Levin's "virtual governor" framing — the idea that a global constraint can be transduced into local incentives so that parts, each minimizing their own discomfort, collectively maintain the whole — for its distributed-control experiments. It is careful, though, to keep this at "scaffold" level: framing vocabulary, not evidence.

2.6 What to carry forward

The agency ideas descend from cybernetics: Ashby's requisite variety and ultrastability (measured in the repo as self-repair), Wiener's feedback (used but under-credited), and Maturana–Varela's autopoiesis (invoked as aspiration, but the strict self-production criterion is not met). Bennett is the modern hub — weakness, the viability buffer, the Law of the Stack, the tapestry of valence — making the program largely an empirical test-bed for his theory, with Levin's multiscale agency supplying newer framing.

CHAPTER 3

Information, Compression, and Constraint

The program's ideas about *generalization* and *geometry* descend from information theory — the mathematics of compression, born in 1948 and central ever since. This chapter traces two strands: the description-length tradition that weakness challenges, and the rate-distortion tradition that gives the program its sharpest (and falsified) prediction.

3.1 Shannon and the bit (1948)

Claude Shannon's *A Mathematical Theory of Communication* founded the field at a stroke, giving the world the **bit** as the unit of information and **entropy** as the measure of a source's unpredictability. The deep idea the program inherits is that *information and compression are the same subject seen from two sides*: the fewer bits you need to describe something, the more structure (and less randomness) it contains. Every later idea in this chapter is a way of asking "what is the right thing to compress, and how?"

3.2 The description-length tradition: Solomonoff, Rissanen, and Occam made precise

The oldest answer to "which explanation should I believe?" is Occam's razor: prefer the simplest. Ray Solomonoff (1964) and Jorma Rissanen (1978, the **Minimum Description Length** principle) made this mathematically exact. Among all models that fit your data, prefer the one that lets you *compress the data most* — shortest (model + data-given-model), measured in bits. A short description is an a-priori probable one; simplicity and compressibility become the same thing. For half a century, this was the dominant formal theory of induction, and it works remarkably well.

In the repo — the rival that weakness dethrones

The program's flagship result is a direct challenge to this tradition. On tasks engineered so that a short "shortcut" rule and a broadly-compatible "weak" rule both fit the training data perfectly, the program pits weakness against an explicit menu of description-length selectors — MDL, compression, shortest-form, plus validation, parameter-norm, and flatness. The finding: the description-length family picks the shortcut and generalizes badly (near 0% of the time), while weakness picks the symmetric rule and generalizes (near 100%). The slogan the program inherits from Bennett to explain this: *simplicity is a property of form, but generalization is a property of function*. A short description commits hard to one guess; a weak hypothesis stays compatible with many, and under distribution shift the weak one travels better. This is the program's most consequential departure from the classical tradition — not a rejection of information theory, but a claim that it was measuring the wrong quantity.

3.3 Rate-distortion and efficient coding: how to spend a limited budget

The second strand asks a different question: given a *fixed* bit budget, how should you spend it? Shannon's **rate-distortion theory** answers in general; the neuroscience of **efficient coding** (from Barlow in 1961 to Ganguli and Simoncelli in 2014) answers for brains. The unifying principle: *allocate resolution in proportion to how much each region matters*. A sensory system with limited neurons should spend them finely where stimuli are common or

important and coarsely elsewhere. W.R. Bennett's 1948 quantization law and its descendants make this a precise power-law relationship between how densely you sample a region and how much error remains there.

Intuition pump — the mapmaker's budget

A mapmaker with a fixed amount of ink must choose where to spend detail. A road atlas renders cities at high resolution and empty deserts at low, because that is where travelers need precision. Rate-distortion theory is the mathematics of that choice — the optimal ink-allocation given a budget and a map of what matters. The program's boldest geometric claim is that a neural code, facing the same budget, should warp its internal "map" the same way: spend representational resolution where *concern* (value, reward) concentrates. "Reward deforms the metric" is efficient coding with a value signal in the driver's seat — the same mathematics the brain's sensory systems are thought to use, pointed at what an agent cares about.

In the repo — a derived law, and its instructive falsification

The program does something admirable: it *derives* the exact form of the reward-deformation law before measuring — resolution should grow as concern raised to the power $d/(d+2)$, giving a specific predicted exponent of 0.5 for a two-dimensional code. Then it measures, and the exponent comes out ≈ 0.30 , not 0.5, with tight confidence intervals. The prediction is *falsified* — but productively: the mismatch implies the code was concentrating its resolution as if it lived in essentially *one* dimension, not two. This is the efficient-coding tradition working exactly as it should: a sharp quantitative prediction, honestly tested, revealing a hidden fact (an effective dimensionality collapse) precisely because it was specific enough to be wrong. It is the program's cleanest inheritance from Shannon's descendants, and its cleanest falsification.

3.4 The through-line: compression under constraint

Step back and the two strands unite in the program's founding thesis (Chapter 1). Generalization (the description-length strand) and geometry (the rate-distortion strand) are both about the same thing: a finite system deciding *what to keep and what to throw away* under a capacity limit. Weakness is a claim about *which* distinctions to keep (the broadly-compatible ones); the reward-deformation law is a claim about *where* to spend resolution keeping them (where concern is high). The concern-weighted-weakness theory paper proves these are two faces of one optimization — which is why information theory, more than any other tradition, is the mathematical spine of the whole program.

3.5 What to carry forward

Information theory casts information and compression as one subject; the description-length tradition (Solomonoff, Rissanen, MDL) made Occam precise and is the rival weakness dethrones with "simplicity is form, generalization is function"; the rate-distortion / efficient-coding tradition (Shannon, Bennett-1948, Ganguli-Simoncelli) says spend resolution by importance, which becomes the program's "reward deforms the metric" law — derived, then productively falsified to an exponent of 0.30; and both strands are the same compression-under-constraint problem, unified in the program's theory.

CHAPTER 4

Symmetry, Topology, and the Brain

The program's most technically distinctive result — that "weakness" is really about symmetry, and that symmetry shows up as a specific *shape* in a neural code — sits at the meeting point of three traditions: the group theory of symmetry, the topology of neural manifolds, and the neuroscience of grid cells. This chapter braids them, because in the program they turn out to be three views of one event.

4.1 Group theory and equivariant machine learning

Symmetry — a change that leaves something looking the same — has been mathematics' deepest organizing idea since the nineteenth century, formalized as **group theory**. In the last decade it transformed machine learning through **equivariant networks**: Cohen and Welling's group-equivariant CNNs (2016), and the "geometric deep learning" synthesis of Bronstein and colleagues (2021), which argued that most successful neural architectures are really exploiting symmetries — a convolutional network works because it is *equivariant* to translation (shift the image, the features shift too). A deep result connects this to Chapter 3's Fourier ideas: the patterns compatible with a symmetry are exactly its **irreducible representations**, which for shift-symmetry are the frequency modes. Symmetry and Fourier analysis are one subject.

In the repo — weakness *is* symmetry-counting

The program's central quantity, weakness, is defined group-theoretically: it counts how many of a symmetry group's transformations a learned function respects. A function compatible with many symmetries treats whole families of inputs alike — it is "weak" in Bennett's sense of committing to little. This is where the program fuses Bennett's abstract weakness with the concrete machinery of equivariant ML: weakness becomes a measurable count, computed against cyclic, dihedral, and permutation groups, and (in the learned-symmetry work) even against a group *inferred from the data itself*. The claim that this count predicts generalization better than any simplicity measure is the program's signature contribution, and it lives squarely in the equivariant-ML tradition.

4.2 Topology and neural manifolds

Topology — the study of shape up to bending and stretching, indifferent to distance and angle — supplies a different, coarser invariant: the count of a shape's "holes" (its Betti numbers). The insight that transformed neuroscience is that the collective activity of many neurons, though embedded in a very high-dimensional space, often lies on a low-dimensional *manifold* with a specific topology. Nina Miolane's work on population-code topology, and the landmark 2022 *Nature* paper by Gardner and colleagues, showed that the brain's grid cells collectively trace out a **torus** — a donut — with the exact hole-signature (1, 2, 1), reproducibly, across animals. Topology gives a substrate-independent fingerprint: two systems can be called "the same" if their activity has the same shape, regardless of what they are made of.

4.3 Grid cells and the convergent code

The grid cells themselves — discovered by the Mosers (Nobel Prize, 2014) — fire in a periodic, hexagonal pattern as an animal moves through space, a biological Fourier basis for position. The result that ties everything together: when you train an artificial recurrent network to do nothing but *integrate its own motion* (path integration), grid cells and their toroidal geometry *spontaneously appear* (Sorscher, Ganguli, and others, 2018–2019). An optimal code for 2-D space, arrived at independently by evolution and by gradient descent — a striking instance of the convergence puzzle of Chapter 1.

In the repo — one event, three needles

The program's grid-cell work makes the deepest unification attempt in the whole enterprise. Its prediction: when a network discovers 2-D translational symmetry, three measurements should all light up together, because they are measuring the same discovery. **Weakness** (the code stays consistent under many shifts), **spectrum** (the code concentrates into a few Fourier modes), and **topology** (those modes trace a torus) become "three measurements of one event: the code discovered the group." It is a genuinely beautiful synthesis linking Bennett's weakness, Fourier analysis, and Gardner's torus. And — in the program's characteristic honesty — the specific *mediation* prediction (that weakness would predict generalization *through* the topology) **failed**: weakness and topology each predict generalization independently, with no single mediating object. The conceptual bridge is real and the honest negative is instructive: here, the paper concludes, "weakness is a footprint, not the cause."

4.4 The Platonic Representation Hypothesis

A 2024 idea from machine learning (Huh and colleagues) crystallized the convergence puzzle for neural networks: the **Platonic Representation Hypothesis** holds that different models, trained on enough data, converge to the *same* internal representational geometry — as if approaching a shared, substrate-independent statistical model of reality. The program adopts this as background and proposes weakness as a *mechanism* for it: when a task has a symmetry, every capable learner is driven to the same high-weakness code, hence the same geometry. This is the program at its most ambitious — offering not just to observe convergence but to explain *when it must happen*.

Verdict — genuine synthesis, honestly bounded

The symmetry–topology–grid braid is where the program is most original: the specific claim that one scalar (weakness) governs whether a code carries the right topological structure of a task's symmetry, tying together three literatures that rarely speak, is a real intellectual contribution. But a mature reader should hold two limits the program itself surfaces. First, the mediation failed — the three needles co-occur but weakness does not *drive* the topology, softening the unification from "cause" to "correlated footprints." Second, the cross-model convergence claim (the Platonic mechanism) did not survive external contact: tested across three unrelated model families, the shared-geometry reading failed. The braid is a beautiful and defensible hypothesis; it is not yet a demonstrated law, and the program's own negatives are what tell you so.

4.5 What to carry forward

Group theory and equivariant ML make weakness a measurable symmetry-count; topology and neural manifolds give the torus as a substrate-independent shape; grid cells show that shape emerging independently in brains and networks; the program braids all three

into "generalization, spectrum, and topology are one event — the code found the group," a genuine synthesis whose specific mediation and cross-model claims honestly failed, leaving a strong hypothesis rather than a proven law.

CHAPTER 5

Causality, Control, and Agency

The program's benchmark and agency work descends from two more traditions: the causal-inference revolution that taught science to distinguish "found" from "made," and the reinforcement-learning and viability-theory lineages that model agents as maintainers of their own existence. This chapter closes Part I by tracing them.

5.1 Pearl and the causal revolution

For most of the twentieth century, statistics was officially agnostic about causation — "correlation is not causation" was a warning, not a method. Judea Pearl's work (culminating in *Causality*, 2009) changed that by giving causation its own mathematics: the **do-operator**, which distinguishes *observing* a variable to have a value from *intervening* to set it. This sounds like a small notational move; it is a revolution. It makes precise why a randomized trial beats an observational study, what "controlling for a confounder" really does, and — crucially for the program — when a hidden quantity is *identifiable* from data at all.

In the repo — the do-operator inside a neural network

The program applies Pearl's framework not to epidemiology but to the insides of models. Its distinction between a feature being *decodable* (a probe can read it) and being *used* (it causally drives the decision) is pure Pearl: decodability is observational, use is interventional. The program tests use by **activation patching** — surgically overwriting an internal feature and measuring the change in the output, which is *do(feature = value)* made literal. And its self/world "gauge problem" — that a system can be a perfect agent while having no fact of the matter about which changes are "its own" — is a causal *identifiability* result: when the data reveals only a sum, the split into parts is underdetermined, fixable only by intervention (the agent deliberately doing nothing). This is Pearl's identifiability theory reincarnated as a philosophy-of-mind result, and it is one of the program's genuinely novel moves.

5.2 Reinforcement learning and its homeostatic turn

Reinforcement learning — agents learning to act by maximizing reward — is the workhorse tradition of modern AI agency. The program uses its standard machinery (policies, value, discounting, the REINFORCE algorithm) but leans on a specific dissenting strand: **homeostatic reinforcement learning** (Keramati and Gutkin), which reframes reward not as an arbitrary scalar to be maximized but as *drive reduction* — the restoration of internal viability variables toward a healthy setpoint. An agent is not a reward-maximizer but a homeostat that stays alive. This reframing is the bridge from the RL tradition to the cybernetic tradition of Chapter 2, and it is why the program's agents have an "energy" that decays and must be replenished rather than a bare score to run up.

Intuition pump — the difference between a gambler and a body

Standard reinforcement learning models an agent like a gambler: maximize the number, whatever it is. Homeostatic RL models it like a body: keep the vital signs in range, and "reward" is just how much a choice moved you back toward health. The two can look similar when the number happens to be "how alive you are" — but they diverge sharply at the edges, where a gambler chases the last point and a body protects its viability. The program deliberately builds bodies, not gamblers, because its whole question is about *agency as self-maintenance*, and you cannot study

self-maintenance in a system that has no self to maintain. This choice — inherited from Keramati, Gutkin, and behind them Ashby — is what makes "concern" a meaningful word in these experiments.

5.3 Viability theory and the free-energy principle

Two more traditions round out the agency picture. Jean-Pierre Aubin's **viability theory** (1990s) formalizes the "viability kernel" — the set of states from which a system can stay alive forever, given its controls and the disturbances it faces. The program uses this vocabulary (its "viability buffer" is the slack between how many futures a policy keeps open and how many it has committed to), though it uses the language informally rather than computing formal kernels. And Karl Friston's **free-energy principle** and **active inference** (2006–present) — the sweeping claim that perception and action both minimize a single quantity, and that action can be *epistemic* (undertaken to reduce uncertainty, not to gain reward) — supplies the program's framing for why an agent should pay to *probe* its own world. The program is careful here: it borrows active inference as a *reading*, a source of framing, not as a formalism it commits to implementing.

5.4 The lineage assembled

Assemble the five chapters of Part I and the program's intellectual ancestry comes into focus. From **cybernetics and autopoiesis** (Chapter 2): agency as self-maintenance, the viability buffer, ultrastable repair. From **information theory** (Chapter 3): weakness as the right compression, the reward-deformation law. From **symmetry and topology** (Chapter 4): weakness as symmetry-counting, the torus as the fingerprint of a discovered group. From **causal inference** (this chapter): the do-operator inside networks, the self/world gauge, availability versus use. And from **reinforcement and viability**: agents as homeostats, not gamblers. Michael Bennett's recent work is the solder joining several of these — weakness, the Law of the Stack, the tapestry of valence — which is why the program reads, at times, as an empirical test-bed for a single contemporary theorist recombining a century of ideas.

Verdict — an honest scholarly gap

One thing a historian should flag: the program's ancestry has unmarked debts. Its core claim that content is grounded in an item's *function for survival* is structurally **teleosemantics** (Millikan, Dretske, Papineau) — a forty-year philosophical literature — yet that tradition is entirely uncited, with meaning-grounding routed instead through Bennett and enactivism. Its "objects form from causal role" is a **natural-kinds** claim (Boyd's homeostatic-property-cluster kinds) that goes unnamed. And its availability-versus-use distinction reinvents Dretske's "structuring versus triggering cause." These gaps make the program look more *sui generis* than it is, and forfeit sharpening tools those literatures already built. Naming them would strengthen, not diminish, the work — a point the philosophy primer develops.

5.5 What to carry forward

Pearl's do-operator and identifiability theory become, in the repo, activation patching and the self/world gauge; homeostatic RL (Keramati–Gutkin) reframes agents as viability-maintaining bodies, not reward-maximizing gamblers; Aubin's viability theory and Friston's active inference supply the buffer and the epistemic-probe framing; Bennett solders much of the lineage together; and the program carries unmarked debts to teleosemantics and natural-kinds theory that naming would strengthen.

PART II

The Trajectory

Part I traced what the program inherited. Part II traces what it did with the inheritance — how a single question fissioned into three research threads, grew a numbered chain of some thirty-five experiments, corrected its own mistakes in public, pivoted when the evidence demanded, and arrived at its current frontier. This is the program as a *living trajectory*, not a static shelf of papers.

CHAPTER 6

The Founding Question and the Three Threads

6.1 The seed

The program begins not with an experiment but with a conceptual wager, stated in its founding synthesis: *"meaning is not merely compression or passive latent geometry. Meaning-like structure appears when differences become salient under concern, and agency-like structure appears when that concern is coupled to action, self-maintenance, boundary preservation, repair, and time."* The whole enterprise is an attempt to take that sentence — which could easily remain armchair philosophy — and make it *tractable*: to find experiments where its claims could be measured, and could be wrong.

Driving every experiment is one recurring anxiety, which is worth naming because it explains the program's entire character: *how can we tell whether a system's behavior depends on the intended internal structure, rather than on a proxy?* A model can get the right answer for the wrong reason. The program's obsession with catching that — its gates, its controls, its "behavior alone doesn't count" rule — all flows from this single worry.

6.2 The fissioning into three threads

The founding question splits naturally into three, according to *what* the geometry is a geometry of:

Thread	The question	Geometry of...
A — Weakness / generalization	Among many rules that fit the training data, which one generalizes to new data — and why?	hypotheses
B — Geometry of concern	Does an agent whose objective is coupled to its viability organize its internal space around what <i>matters</i> , and does that organization causally drive behavior?	representations
C — Agency / benchmark	Can a minimal agent predict its viability, intervene on its world, attribute changes to self versus world, and maintain that boundary over time?	action & attribution

These are not three separate projects but three faces of one hypothesis. Thread A asks which distinctions a system should keep (the weak, broadly-compatible ones). Thread B asks whether a living system keeps the ones that *matter* to it (concern-weighted). Thread C asks whether keeping them adds up to *agency* (self-maintenance under intervention and time). Much of the program's later theoretical work is the attempt to prove these three threads are mathematically one — that weakness, concern, and agency are views of a single optimization. Part I's Chapter 3 and the math primer's theorems chapter are where that unification is attempted.

Intuition pump — three instruments, one phenomenon

Imagine astronomers studying one object with three instruments — a telescope, a spectrometer, a radio dish. Each sees something different (an image, a chemical signature, a pulse), and the science advances by showing these are one star. The program's three threads are three instruments trained on one conjecture about meaning and agency.

Thread A's weakness-counting, Thread B's cluster-gap geometry, Thread C's intervention-and-attribution — each measures a different facet, and the program's ambition is to demonstrate they are reading one underlying thing. Whether they truly are, or only seem to, is the open question the whole trajectory circles.

6.3 The character this gives the program

Two features of the program's style follow directly from this founding structure. First, its **relentless anti-proxy discipline**: because the founding worry is "did it work for the right reason?", every result must survive controls that a merely-behavioral success would fail. Second, its **layered, ladder-like growth**: because the three threads are meant to be one, each result tends to expose the next question, producing a chain rather than a scatter. That chain — the "numbered spine" of some thirty-five experiments, each answering the previous one's residual doubt — is the subject of the next chapter, and it is the most distinctive thing about how this program grew.

6.4 What to carry forward

The program began as a wager that meaning is "geometry under concern" and agency is concern coupled to self-maintenance; it is driven throughout by one worry — did the system succeed for the intended reason or a proxy?; the wager fissions into three threads (weakness/generalization, geometry-of-concern, agency/benchmark) meant to be mathematically one; and this founding structure produces the program's anti-proxy discipline and its ladder-like, self-correcting growth.

CHAPTER 7

The Numbered Spine and the Correction Chain

7.1 A program that grows like a proof

Most research programs grow by branching — many groups pursuing loosely related questions. This one grew like a *proof*: a numbered sequence of experiments (Papers 1 through 25 and beyond), each one triggered by a doubt the previous one could not dispel. The program calls the underlying pattern its **Correction Chain**, and it is the single most distinctive thing about its history. Each link has the same shape: *a system succeeds behaviorally* → *a diagnostic reveals it succeeded via a proxy, not the intended structure* → *a small architectural change makes the structure load-bearing* → *a harder setting reveals the next proxy*. Repeat.

Intuition pump — peeling an onion of self-deception

Each layer of the onion is a way the system could *look* right while being wrong, and each experiment peels one off. The agent survives — but is it representing its viability, or just reacting to a sensory cue? Peel: it turns out to use a proxy. Add structure so it must represent viability. Now it represents viability — but can it tell which changes it caused versus the world caused? Peel: the two are confounded. Add an intervention to separate them. Now it attributes correctly — but does it keep attributing after the world changes? Peel again. The program's history is this peeling, and its honesty is that it keeps finding new layers rather than declaring the onion finished. Reading the numbered papers in order is watching a system refuse to let itself off the hook.

7.2 The chain, in brief

The spine is long; here is its logic in compressed form, so you can see the doubt propagating from each result to the next:

Stage	Result	The doubt it raised
Weakness (P1–3)	Weakness predicts generalization where simplicity fails; the symmetry can even be inferred from data.	But on real language models the geometry is <i>passive</i> — it predicts, but does it <i>control</i> ?
Passive → active (P4)	Coupling a representation to action makes it causally load-bearing (a 7× jump).	Is that real self-maintenance, or just "of course ablating a trained direction hurts"?
Autopoietic (P5)	The active representation also repairs itself after damage and obeys the Law of the Stack.	Do <i>objects</i> themselves form from concern, or only directions?
Valence objects (P6)	A viability-coupled encoder carves the world by causal role, discarding sensory similarity (reward-axis +1.96 vs sensory +0.005).	Can this representation actually drive <i>competent action</i> ?

Two bottlenecks (P7–9)	Representation and competence are independent: a good geometry can fail under weak learning signal, and greedy planning on a learned model <i>is</i> a competent policy.	Is the training distribution even stable enough to learn from?
Boundaries & uncertainty (P12–14)	Smooth networks fail at sharp decision boundaries; ensembles' "uncertainty" does not track error.	Can a single drive even represent multi-dimensional concern?
Vector self (P15–16)	Vector-valued concern enables priority reweighting; but the self/world split is <i>gauge-symmetric</i> — right behavior, arbitrary attribution.	Can the agent choose <i>when</i> to probe to break the gauge?
The probing saga (P17–21)	A long sub-chain: naive probing saturates → a fix is anti-calibrated → recomputing against the current model fixes it → the same failure recurs on a second dimension → scale-normalization fixes that too.	What if the world doesn't hold still while you learn?
Responsive world (P22–25)	When the world responds to the agent and shifts, the agent self-silences, then over-probes ("anxiety"), then achieves a stable detect-probe-cool-reopen cycle — but a shared network can't disambiguate role-specific effects.	An <i>architectural</i> ceiling: the next question is structural, not about training.

7.3 The probing saga: the chain at its most granular

One sub-chain deserves a closer look because it shows the correction discipline at maximum resolution. The question was simple: an agent should probe (run a costly experiment on its own world) where it is *uncertain*. The obvious rule — "probe where your prediction error is high" — was tested and *falsified*: high error can be irreducible noise you cannot learn from, so the naive prober fired constantly. The fix (a running average of past errors) turned out *anti-calibrated* — it pointed reliably at the *wrong* places. The next fix recomputed the error against the *current* model rather than a stale average, and finally flipped the calibration positive. Then the identical failure reappeared on a second internal dimension, requiring a scale-normalization fix. Four papers, each correcting the last, to get one signal right — and each correction published, with the failed intermediate versions preserved rather than hidden. This is the program's method in miniature: not a leap to the answer, but a documented descent through the ways of being wrong.

Verdict — the strength and the risk of a proof-like program

The correction chain is genuinely admirable: it is what serious self-scrutiny looks like, and it produces real knowledge (the value-of-information insight, the gauge result, the architectural ceiling are all durable). But a mature reader should hold one risk. A program that grows as a single chain, each link authored and judged by the same system, can descend a long way into a *self-consistent* structure without an external check that the whole ladder is leaning against the right wall. Each correction is validated against the program's own environments and gates. The peeling is real; whether the onion is the right onion — whether these toy worlds teach us about real minds and models — is a question the chain cannot answer from inside itself. That is why the program's later turn toward external models and data (Chapter 8) matters so much.

7.4 What to carry forward

The program grew like a proof — a numbered spine of ~35 experiments linked by a Correction Chain (behavior succeeds → proxy exposed → structure added → next proxy revealed); the chain runs from weakness through passive-to-active geometry, valence objects, the two-bottlenecks dissociation, the gauge problem, the four-paper probing saga, to a responsive-world architectural ceiling; its granular self-correction is its great strength; and its risk is that a single self-judged chain can grow self-consistent without an external check on whether its toy worlds generalize.

CHAPTER 8

Arc 2, the Benchmark, and the Phases

8.1 The director's reset

Around the midpoint of its history, the program did something research programs rarely do to themselves: it delivered a blunt internal critique and changed course. The founding numbered arc, the critique held, had become "too demo-shaped" — its experiments tested whether an agent could *use a provided probe well*, not whether it could *invent the next useful experiment*. Passing a supplied test is a lower bar than deciding what test to run. This reset, visible in the program's phase documents, pushed it from "mechanism" toward what it calls "regime transition" — results that change what can even be represented.

Intuition pump — from answering the exam to writing it

There is a difference between a student who aces every exam you hand them and a student who figures out which experiment the field should run next. The first has mastered a fixed game; the second is doing science. The program's Arc 2 reset is the recognition that its agents had been acing handed exams, and the decision to test something harder: can an agent *invent the intervention* — choose, from a space of possible experiments, the one that would reveal a hidden structure it needs? Inventing the right question is much harder to fake by pattern-matching than answering a given one, which is precisely why the program pivoted toward it.

8.2 Arc 2A and 2B: from using probes to inventing them

Arc 2A ("concerned syntax") built agents that must, from raw pixels, *choose* which intervention to run to expose a hidden compositional structure — and must do so only when that structure actually matters for their survival ("concern-gated intervention invention"). It climbs a ladder of increasing autonomy: from a supplied vocabulary of experiments, to label-free calibration, to discovering the relevant categories itself. **Arc 2B ("viable computational bodies")** went a level up: it evolves candidate agent *architectures* — "bodies" — under viability and formal constraints, using a typed Haskell checker to gate which bodies are even admissible. The two arcs couple at a shared contract, with a "scaffold-removal ladder" that progressively strips away supervision to test how much the agent can do on its own. Both reach their transfer targets under anti-cheat controls — while honestly leaving natural-image perception and open-ended semantics out of reach.

8.3 The benchmark: packaging the science for the field

The program's next move was to turn its accumulated agency findings into a *product* the wider research community could use: the **Causally Grounded Finite Agents Benchmark**. Its thesis, distilled from the whole trajectory: an agent should be judged not on whether it gets answers right but on whether it succeeds *for the right causal reasons* — whether "the variables that matter for future control are represented, attributed, queried, and committed at the causal surfaces where action is selected." Its enforced minimum: no suite passes on behavior alone; a pass requires behavior plus at least one structure-specific gate plus anti-cheat controls.

The benchmark crystallized the trajectory's hard-won lessons into four named "laws of machine agency" — compact statements of things the correction chain had discovered the hard way:

Law	What the chain learned
Predictive Policy Closure	An agent that learns how its actions change what it cares about can use that consequence model directly as its policy.
Reafferent Identifiability	Splitting inputs into "self" and "world" is not identification; telling them apart requires a gauge-breaking signal (like a deliberate null action).
Re-Engagement Floor	Efficient inquiry is not adaptive inquiry; quiet is not evidence of stability. An agent must re-probe when the world changes.
Commitment-Surface Memory	Memory becomes agent-relevant precisely where it is coupled to a future decision.

8.4 The phases: leaving the toy world

Running alongside the benchmark packaging, a sequence of "phases" (4, 5, 6) addressed the deepest worry from Chapter 7: *does any of this hold outside the worlds the program built itself?* The director stated the critique bluntly — "every result lives on toy worlds the lab built itself" — and the phases were the response. Phase 4 tested which mechanisms could be learned at all; Phase 5 pushed them toward "more model-like" settings; Phase 6 ran them on *actual public models* (small language models, frozen sentence encoders). This is the program at its most exposed, and the results were mixed in exactly the way honest external contact produces: some mechanisms survived (uncertainty decoupling from error, value-weighted metric deformation on real encoders), while the flagship weakness → OOD prediction was **"hard-killed"** on an external model — its correlation with generalization essentially zero.

Verdict — the phases are the program's moment of truth

A mature reader should treat the phases as the most important part of the trajectory, because they are where the program's claims meet ground it did not prepare. The honest scorecard: the agency and uncertainty mechanisms show real, if bounded, external signal; the concern-geometry story does *not* yet transfer to language as-is (a wrong-sign result); and the flagship weakness result may be bound to self-built worlds — a single external hard-kill that forced the program's largest theoretical reframe (Chapter 9). That the program ran these tests at all, and published the kills, is to its enormous credit. That most of its boldest claims have not yet survived external contact is the plain state of the evidence, and the program says so.

8.5 What to carry forward

A mid-program reset pushed from "using probes" to "inventing them"; Arc 2A/2B built agents that choose interventions and evolve their own architectures under a typed formal gate; the work was packaged into the Causally Grounded Finite Agents Benchmark ("succeed for the right causal reasons") with four named agency laws; and Phases 4–6 carried the claims to real models, where agency and uncertainty mechanisms showed bounded external signal but the flagship weakness result was hard-killed — the program's honest moment of truth.

CHAPTER 9

The Reframe and the Frontier

9.1 When an anomaly reorganizes a program

The history of science has a recurring shape: a program accumulates results within a framework until a single stubborn anomaly forces the framework itself to change. This program had such a moment. The external "hard-kill" of Chapter 8 — weakness failing to predict generalization on a model the lab did not build — was not just a failed experiment; it was a crack in the foundation. If weakness is the *cause* of generalization, it should not vanish when you change the model. Its vanishing implied that weakness might be, in the program's own later phrase, *a footprint rather than the cause*.

9.2 The commitment-surface reframe

The response is the program's current organizing theory, and its title states the reframe: "*Weakness Is a Footprint, Compatibility Is the Cause*." The move is subtle and worth understanding, because it reorganizes the entire prior trajectory. The program had, throughout its history, been treating *availability* of a structure — a probe can decode it, weakness predicts it on a toy world — as evidence the structure was *real* and used. The reframe demotes this. A structure is real for a deployment only if a *train-time intervention* that couples the model to the actual deployment-generating process produces a *causal effect at the commitment surface* — the place where the decision is actually made — and only if that effect survives gauge-fixing and changes of context. Weakness and geometry become *diagnostics* — footprints of the real cause — load-bearing only when the symmetry you score against happens to match the one that actually generates the deployment data.

Intuition pump — the footprint and the animal

A footprint in mud tells you an animal passed. It is real evidence — but the footprint is not the animal, and in the wrong terrain (rock, water) the same animal leaves no print at all. The program spent years measuring footprints (weakness, cluster geometry, probe decodability) and treating a clear print as proof of the beast. The reframe insists on tracking the animal itself: the causal effect at the point of decision. Footprints remain useful — they are cheap and often reliable — but only where the ground takes prints, i.e. only where the measurement's assumptions match reality. The external hard-kill was a patch of rock where the animal left no print, and it forced the program to stop mistaking prints for beasts. This is a genuine conceptual maturation, driven by a negative result the program could have buried and did not.

9.3 A confound the program found in its own flagship

The reframe also surfaced a subtle confound in the program's best result, and the way it handled this is characteristic. Its data-augmentation scheme, it realized, placed *correctly-labeled* examples onto the held-out deployment region — meaning a model might score well not by *learning the underlying rule* but by effectively *memorizing transformed copies of labels*. The existing experiments could not tell these apart. Rather than let the ambiguity ride, the program pre-registered a new multi-arm experiment (the "E5" generator-versus-coverage design) built specifically to separate "learning the rule" from "labeled coverage." At the time of writing, that experiment is at the frontier — frozen, partly

run, verdict pending. The program is, in effect, actively trying to determine whether its own flagship result is real or a sophisticated artifact.

9.4 The current frontier: from verifying to learning

The frontier has a clear theme: the crossing from *static verification* (does this trained system use the structure?) into *learning dynamics* (does the commitment-surface criterion govern how a system *improves itself* over time?). Several experiments sit pre-registered and mostly unbuilt, each testing whether the program's hard-won criterion is not just a diagnostic but a usable *objective*:

- **Self-training without collapse.** The flagship frontier test asks whether a commitment-surface-survival *reward* can prevent the reward-hacking collapse that self-training loops are known to suffer — whether "does the structure reach the decision surface?" is a self-improvement objective, not merely a verification check.
- **Continual learning by protecting the load-bearing subspace.** Whether protecting only the specific compatibility-carrying directions (rather than all weights uniformly) escapes the usual stability-plasticity tradeoff.
- **Commitment-change as a plasticity trigger.** Whether the moment a commitment changes is a better signal for when to adapt than any internal statistic.

These share a hypothesis: that the criterion the program spent its whole history discovering — a structure is real when it causally reaches the commitment surface — might be the right thing not just to *measure* but to *optimize*. If so, the program crosses from a science of diagnosing agency into a science of building it. If not, the criterion remains a valuable measuring stick. Either way, by the program's own design, the result narrows the program.

Verdict — a program that reframes on its own negatives is doing it right

The commitment-surface reframe is the strongest single piece of evidence that this is a real research program and not an elaborate confirmation machine. A confirmation machine does not demote its own flagship quantity to "a footprint" on the strength of one external kill; it explains the kill away. This program reorganized its entire theory around the anomaly, then designed experiments to test whether its best result is an artifact. That is exactly the behavior philosophy of science prescribes and rarely observes. The open question — whether the new criterion survives its own frontier tests — is genuinely open, which is the honest and correct state for a frontier to be in.

9.5 What to carry forward

An external hard-kill cracked the "weakness is the cause" framework; the commitment-surface reframe ("weakness is a footprint, compatibility is the cause") demoted availability in favor of causal effect at the decision surface; the program found and is actively testing a confound in its own flagship result; and the frontier is the crossing from verifying that a system uses structure to using the criterion as a self-improvement objective — with the whole reframe standing as proof the program corrects itself on its own negatives.

CHAPTER 10

The Meta-Story: Machine-Made Science

10.1 A history unlike other histories

Every previous chapter could, in principle, describe a program run by human scientists. This one cannot, because the most historically novel fact about this program is *how* it was made: its experiments, sweeps, figures, and paper drafts were **generated by an AI coding agent under a human director's** (Jawaun Brown) direction and review. This is not a footnote to the history; in an important sense it *is* the history, because it shapes the program's pace, structure, and character in ways worth reading explicitly.

10.2 How the machine-authorship shows up

Feature	How AI-authorship produced it
Pace and volume	Roughly fifty experiments and as many papers in a span no small human team could match — a numbered arc, three benchmark suites, and multiple phases, all inside one compressed period.
Self-documentation as infrastructure	An unusually heavy scaffolding of ledgers, audit cards, provenance files, and handoff documents <i>written for other AI agents</i> — because each agent may start with no memory of the last, so the repository itself must be the memory.
Guardrails against machine failure modes	The pre-registration apparatus, anti-cheat gates, publication guards, and "do not retune a gate" rules read as defenses specifically against the ways an eager generator would overclaim: retroactive reframing, metric-gaming, leaking answers into the setup.
Rhetorical restraint as policy	The near-verbatim "we do not claim consciousness" disclaimers in paper after paper are a <i>systemic</i> guard against a vocabulary (concern, self, care) that constantly invites overreach.

Intuition pump — a laboratory built for amnesiac scientists

Imagine a lab staffed by brilliant researchers who each forget everything at the end of the day and are replaced by a fresh colleague every morning. To function, such a lab would have to externalize all its memory into meticulous notebooks, standing procedures, and handoff notes — because nothing lives in anyone's head overnight. This program is that lab. Its extraordinary density of ledgers, provenance cards, and "next-agent handoff" documents is not bureaucratic excess; it is the necessary condition for a research program whose workers are stateless. The repository is not just where the science is stored; it is the only place the science's continuity *exists*. This is a genuinely new organizational form, and its documentation habits are the fossil record of it.

10.3 The human in the loop

Where, in all this, is the human director? Most visible not in the experiments but in the *taste and ambition*. The program's sharpest internal documents — the mid-program critique that forced the Arc 2 reset, the demand to leave the toy worlds, the judgment that most results are "only diagnostic-to-mechanism" — read as a human steering the level of ambition while the agents supply throughput. The director judges the program through the lenses of named

researchers (the compute-scaling view, the causal-inference view, the "build the strongest shortcut first" ethic) and pushes back when the output drifts toward "one more positive table" instead of "change the allowed claim boundary." The division of labor is legible: the machine generates and self-checks at scale; the human sets the taste, picks the fights, and decides what would actually count as progress.

10.4 The novelty question, answered honestly

A history owes its reader a verdict on originality. What, across this whole trajectory, is genuinely new versus skilled recombination of the lineage in Part I?

Verdict — what is genuinely the program's own

Defensibly novel: (1) turning Bennett's philosophical "weakness" into a reparameterization-invariant, data-inferable, *causal* predictor of generalization that beats every classical simplicity measure in a clean head-to-head — the crown jewel. (2) The passive → active transition made *measurable*: showing that coupling a representation to action turns a correlational cluster into a self-defending, self-repairing attractor. (3) The self/world *gauge* result — a computational proof that a system can be a perfect agent while having no fact of the matter about which changes are its own, fixable only by intervention. (4) The "availability is not use" / load-bearing standard as a general *criterion of representational reality* — which survives all of the program's own negatives and may be its most durable contribution. (5) An unusually large corpus of *honest negatives* — the falsified exponent, the failed mediation, the external hard-kill, the self-retracted synthesis — published at a density that is itself a contribution to how AI-generated science should be done.

Skilled recombination (discount accordingly): the grand convergence thesis (a synthesis of the efficient-coding, Platonic-hypothesis, and enactivist observations); the "concern / viability / boundary" vocabulary (recombined from Ashby, Maturana–Varela, Di Paolo, Bennett, Levin); and most of the mathematical machinery (equivariance, rate-distortion, viability kernels, causal identifiability), imported largely intact.

10.5 The bottom line

Read as a whole, the trajectory is a program that took a cluster of speculative philosophy-of-mind claims and did the unglamorous, disciplined work of trying to make them measurable and to catch itself when they weren't. Its strongest results sit at the intersection of generalization, geometry, and intervention on small clean systems; its reach exceeds its grasp when it stretches toward real models and open-ended meaning, and it mostly says so. The historically new thing is not any single result but the *form*: a fast, self-documenting, self-critiquing, honestly-negative research program authored by a machine and steered by a human — an early specimen of what science may look like when its labor is automated but its judgment is not. Whether its boldest bets are right remains open. That it made the bets legible, testable, and honestly scored is already a contribution, and it is the one a history can certify.

10.6 What to carry forward

The program's most novel fact is its authorship — AI-generated under human direction — which explains its pace, its memory-as-infrastructure documentation, its guardrails against machine overclaiming, and its rhetorical restraint; the human supplies taste and ambition, the machine supplies throughput and self-checking; genuinely novel contributions include causal weakness, measured passive → active transition, the self/world gauge, the load-bearing criterion, and the density of honest negatives; and the deepest novelty is the *form* — an early specimen of automated-labor, human-judgment science.