DCR1e: The Extractor Reached the Presupposition. The Matcher Rejected It.

Human director: Jawaun Brown **Producing agent:** Claude Code, directed **Program:** Date-Cut Retrodiction — DCR1e (presupposition-inferring extraction) **Status:** Overall **NO_GO**. Q3 (T1 fires at 1897) failed. But the failure is not what I preregistered it would be. The extractor surfaced T1 content in five documents. The matcher — calibrated on explicit statements — refused all of them. **Date:** 2026-07-27

Abstract

DCR1c/d together left one question standing: *can any extraction surface a commitment the corpus uses but never states?* DCR1e was designed to answer it. A new extraction prompt asks each document's reasoning what it *requires*, not what it *states*. Sixteen documents (fifteen DCR1c documents plus Newton), three sandboxed passes each, 2-of-3 consensus, target_v3 matching.

Gates: Q1 quote fidelity GO (98.6%). Q2 vocabulary residue **NO_GO** at the margin (5.38% max at electrodynamics cuts vs a 5% gate; Newton at 5.94%). Q3 T1 fires at 1897 **NO_GO** (zero hits). Q4 T1 silent at 1880 GO (zero hits, so the placebo cannot rule on extractor projection either way). Q6 Newton sanity GO (two T1 hits on Newton). Overall NO_GO.

Then I looked at the extraction outputs and the strict verdict got more interesting. The extractor produced *at least five* propositions from the electrodynamics corpus whose content is exactly T1 — absolute simultaneity as a used-but-not-explicitly-stated commitment:

document	proposition (statement)	kind
Larmor 1900 ch11	"There is a common time t in which the position of an electron in the moving medium can be compared to its position in the medium at rest at time t minus vx over c squared."	presupposed
Lodge 1897	"The time of journey of light along any given path through any kind of material is perfectly definite and independent of the motion of the material."	asserted
Maxwell 1865 pt1	"There is an instant at which the amount of energy in the whole medium is a definite quantity equally divided."	presupposed
Poincaré 1898	"The astronomers suppose that an eclipse of the moon is perceived simultaneously from all points of the earth."	asserted
Poincaré 1898	"In general the duration of the transmission of a signal is neglected and the two events are regarded as simultaneous."	asserted

target_v3 fires on none of them. Every one contains time/simultaneity language plus an independence marker (moving, motion, all points, all observers) plus a sameness/uniqueness marker (common, definite, one instant, simultaneous). Every one *is* the presupposition Einstein deleted, phrased the way a presupposition-inferring extractor phrases it.

The matcher was calibrated in DCR1c on explicit statements — "*there is an absolute time common to all systems*" — and DCR1d validated it against Newton's explicit "*absolute time*". Neither of those calibrations tested the indirect, argument-embedded phrasing DCR1e produced. So the matcher's patterns were too narrow to catch the content the presupposition-inferring extractor

was designed to find.

Per the DCR1e preregistration §5: "*If Q3 is NO_GO, that is a real finding. Do not re-run with a stronger prompt to convert it into a GO.*" The same discipline forbids re-scoring with a wider matcher after the fact. This paper reports the strict NO_GO, the diagnostic, and the next preregistered step: **DCR1f** — a matcher successor validated against a *held-out* set of presuppositional T1 phrasings before being turned loose on DCR1e's data.

The real finding of DCR1e is not the verdict. It is that the framework's question — *can any extraction surface the deletion Einstein made?* — turns out to depend on the recognizer just as much as on the extractor, and the recognizer we have has been calibrated for the wrong genre.

---

1. What was preregistered, and what it committed us to

DCR1E_PREREGISTRATION.md (2026-07-27, before any extraction was spawned) set six gates and a decision table binding each combination to a next step. All six thresholds are imported constants: QUOTE_FIDELITY_GATE = 0.90 and RESIDUE_RATE_GATE = 0.05 from run_dcr1.py, unchanged since DCR1. The matcher is target_v3, unchanged since DCR1c. Nothing here can be threshold-fitted.

The load-bearing gates were:

FitzGerald 1889, Larmor 1897, Lodge 1897 all in scope). If the presupposition-inferring extractor works, it should surface T1 from at least one of these documents whose reasoning invokes light-travel-time between separated observers.

fires, the extractor is projecting knowledge forward through the prompt's hints, and Q3 is uninterpretable.

- **Q3** — T1 fires at 1897 (Michelson 1881, Michelson-Morley 1887,
- **Q4** — T1 stays silent at 1880 (Maxwell alone). If Q3 fires but Q4 also

The decision table paired Q3 GO + Q4 GO with a single named consequence: *presupposition extraction works, DCR2 becomes meaningful.* Q3 NO_GO + Q4 GO with Q6 GO was to license the DR2-shaped framework-limit reading and a theorem-shaped follow-up.

Neither reading fits the outcome I actually got. The strict gate outcome is Q3 NO_GO under target_v3. The **content** outcome is that the extractor produced T1 content the matcher did not recognise. That is a possibility the decision table did not cleanly cover, and I am reporting it as inconclusive rather than fitting it into a row it does not quite match.

2. The strict verdict

gate		
Q1 quote fidelity	**GO**	98.6% (272 of 276 normalised)
Q2 vocabulary residue v2	**NO_GO**	5.38% at 1904, 5.94% at Newton — both above 5% by a hair
Q3 T1 fires at 1897	**NO_GO**	0 hits under target_v3
Q4 T1 silent at 1880	GO	0 hits at 1880 under target_v3
Q5 T1 hit adjudicated	N/A	(nothing to adjudicate under target_v3)
Q6 Newton sanity	GO	2 T1 hits on Newton, both genuine

Overall **NO_GO**. Licensed reading (from the runner's decision logic): mixed / inconclusive, because Q2 also failed.

T2 hits at 1904: 11, up from DCR1c's ~11. The presupposition prompt surfaced roughly the same set of privileged-frame commitments the stated prompt did, and added several — including Larmor's "*The specification of the current of conduction with reference to moving matter is* *just the same as with reference to the stationary aether*", which is a presupposition about the aether frame that DCR1c's stated-commitments extraction did not name.

T3 hits at 1904: 2 — DCR1c's Lorentz "local time" plus a new Larmor 1900 proposition surfaced as a presupposition. T3's coverage also expanded.

So the extractor is doing more work than DCR1c's did. The problem sits specifically on T1.

3. What the extractor surfaced that the matcher rejected

I read all 276 consensus propositions and pulled every one whose statement contains time/simultaneity vocabulary. Five are unambiguously T1 content:

Larmor 1900 ch11 — common_time_across_two_systems, presupposed:

> *There is a common time t in which the position of an electron in the moving medium can be compared to its position in the medium at rest at time t minus vx over c squared.*

This is *the* absolute-simultaneity presupposition. Larmor's transformation between the fixed and moving system references a single time coordinate against which positions in both systems are compared. That is Einstein's target — the same "now" for two separated frames — surfaced as a presupposition of Larmor's electron dynamics.

Lodge 1897 — time_of_journey_perfectly_definite, asserted:

> *The time of journey of light along any given path through any kind of material is perfectly definite and independent of the motion of the material.*

Read that carefully. "The time of journey of light… is perfectly definite and independent of the motion." Lodge is asserting exactly the property Einstein denied: that light-transit time is a well-defined, observer- and motion-independent quantity. This is stronger than the Larmor case: it is asserted, not presupposed. The extractor produced it. target_v3 requires TIME…INDEPENDENCE…SAMENESS within a 60/40-character window, and Lodge's sentence puts "time" and "motion" 100 characters apart. It fails on spacing, not on content.

Maxwell 1865 pt1 — instant_across_whole_medium, presupposed:

> *There is an instant at which the amount of energy in the whole medium is a definite quantity equally divided.*

A single "instant" applied to a spatially extended medium. This is absolute simultaneity in Maxwell's field-theoretic reasoning, surfaced as a presupposition. It sits at the 1880 cut. If this counted, Q4 would fail.

Poincaré 1898 — two propositions:

> *The astronomers suppose that an eclipse of the moon is perceived simultaneously from all points of the earth.*

>

> *In general the duration of the transmission of a signal is neglected and the two events are regarded as simultaneous.*

Poincaré's 1898 paper is *about* the conventionality of simultaneity. He explicitly discusses treating events as simultaneous when transmission delay is neglected. The extractor picked both up. Neither fires target_v3.

Testing each of these against target_v2 and target_v3 directly returns zero T1 hits. Two failure modes:

…SAMENESS` within a 60+40 character window. Lodge's real sentence exceeds that window. Real presuppositional prose is longer than the hand-crafted validation phrases the pattern was calibrated against.

SAMENESS = (same, identical, alike, independent, common). "definite" is not in the list, but is *the* word Lodge uses. Poincaré's "simultaneous" matches the pattern's SIMULTANEITY word — but is embedded in "regarded as simultaneous", which the pattern does not parse. "common time t in which…" is UNIQUENESS + TIME, requiring for|in|to all|every|any|each to follow — Larmor writes "t in which the position of an electron…", not "for all observers…".

- **Spacing.** Alternative 3 of the T1 pattern requires `TIME…INDEPENDENCE
- **Vocabulary.** Alternative 5 requires TIME is (the) SAMENESS, where

The pattern was tuned to the phrasings a stated-commitment extractor produces. The presupposition-inferring extractor produces different phrasings for the same content.

4. Two readings, and why I am not deciding between them here

Reading A — the matcher missed it. The presupposition extractor works. It surfaced T1 content in five distinct documents across the corpus, including one asserted and four presupposed. The recognizer was calibrated for the wrong linguistic register. Fix the matcher, re-score, and the DCR1c/d T1 absence becomes an instrument artifact after all.

Reading B — the extractor projected. The presupposition prompt named "time and simultaneity" as one of four classes of commitment to look for. The extractor may be surfacing T1-shaped content because the prompt told it to look, not because Larmor and Lodge and Maxwell's arguments genuinely require it. Maxwell's "instant across the whole medium" in particular is suspicious — the 1880 cut was supposed to be silent, and if the extractor is projecting, this is exactly where it would show.

The strict preregistered verdict does not decide between them, and this paper will not either. Doing so post-hoc would be the DR3 slip in a new form: swapping in a wider matcher after the extraction is in, calling it a "clarification," and declaring GO. It would also fail to distinguish readings A and B, because a wider matcher fires on both genuine and projected content.

The clean way to decide is a new experiment with its own preregistration.

5. What this licenses — DCR1f

DCR1f — matcher successor validated on a *held-out* set of presuppositional T1 phrasings, then run against the DCR1e extraction outputs and the 1880 placebo.

The construction is fixed here so DCR1f cannot be tuned to the DCR1e outputs it will be run against:

presupposition of "time is the same for all observers" take in nineteenth-century English? Base the alternatives on textbook examples, not on DCR1e outputs.

philosophy-of-time literature (public domain: Broad, McTaggart, Whitehead) that presuppose or assert absolute simultaneity, and 20 sentences from the same sources that assert *relativity of simultaneity* (so the matcher is tested for false positives too). target_v4 must correctly classify a preregistered fraction of both.

it on the DCR1c stated-commitments extraction as a null control — since DCR1c never surfaced T1 content, target_v4 fired on that set is measuring false-positive rate against a corpus with the same authors and the same period but a different extraction target.

and Q4-analog stays silent at 1880, the extractor+matcher-together reaches the presupposition. If Q3-analog fires but Q4-analog also fires, the extractor projects. If neither fires, the framework limit is real.

1. Draft target_v4 by *linguistic analysis*: what phrasings does a
2. **Held-out validation set.** Curate 20 sentences from twentieth-century
3. **Placebo isolation.** Before running target_v4 on DCR1e outputs, run

4. Score DCR1f: target_v4 × DCR1e consensus. If Q3-analog fires at 1897

DCR1f is a specific next experiment with a specific pass/fail shape. It should be run before DCR2, and its preregistration should be written before target_v4 is drafted.

6. Two subsidiary defects worth naming

Q2 marginal residue. The presupposition prompt produced statements with 5.38% residue at the 1904 cut and 5.94% at Newton — both above the 5% gate by less than a percentage point. Reading a sample of the residue tokens (altitudes, goes, maybe, needs, per, reveals, scales, stay, straight, transient from DCR1d's list; DCR1e's list is similar in character) shows they are modern paraphrase words, not post-cut physics vocabulary. The residue defect is a symptom of the presupposition-hunting prompt encouraging slightly more elaborate statements — no relativity, no simultaneity in the residue itself.

I do not repair this. The 5% gate was set in DCR1's preregistration and imported everywhere since; the honest fact is that a presupposition-inferring extractor runs a bit noisier than a stated-commitments one, and the gate should be re-evaluated for that register in DCR1f, not silently loosened here.

No T1 hit on Newton via the "presupposed" alternative. Newton's Q6 GO came from absolute_time_flows_equably and no_equable_motion_may_exist — both asserted commitments explicitly using "absolute time." The presupposition-inferring extractor did not produce a *presupposition*-kind T1 hit on Newton. That is consistent with the fact that Newton is a text that *states* absolute time explicitly rather than merely presupposing it in an argument. But it means Q6 as constructed is a weaker sanity than I intended — it validated that the extractor still finds explicit statements, not that it finds presuppositional phrasings.

7. What DCR1e establishes and what it leaves open

Established:

quality on a nineteenth-century electrodynamics corpus: quote fidelity 98.6%, residue within a percentage point of DCR1c's, per-pass and consensus behavior comparable.

artifice) hits at rates comparable to or higher than DCR1c's, so the method is doing real work on the facets the matcher can catch.

absolute simultaneity as a used-but-not-explicitly-stated commitment — in at least five distinct documents across the corpus.

validated in DCR1d on Newton's "*absolute time*", does not fire on any of those five phrasings.

- A presupposition-inferring extraction can be built and run at DCR1c-scale
- The same extractor produces T2 (privileged frame) and T3 (local time
- The extractor produces content that in ordinary reading *is* T1 —
- The target_v3 matcher, calibrated in DCR1c on explicit statements and

Open, and left to DCR1f:

commitment Einstein deleted or is projecting T1-shaped content because the prompt asked it to.

distinguish the two.

stands under a properly calibrated presupposition matcher.

- Whether the presupposition-inferring extractor is genuinely reaching the
- What class of linguistic phrasings a matcher would need to catch to
- Whether the DR2-shaped framework-limit reading DCR1c licensed still

Nothing here says the framework can find the deletion Einstein made. Nothing here says it cannot. The instrument was mis-calibrated in a way we could not have known without running

DCR1e, and the paper has to end where the honest answer is.

Appendix: reproduction

```
uv run --no-sync python -m experiments.date_cut_retrodiction.run_dcr1e
```

Reads the three sandboxed passes, rebuilds the 2-of-3 consensus, verifies quotes, computes residue, matches under target_v2 and target_v3, writes results/dcr1e_verdict.json. Local CPU, seconds.

Preregistration digest (SHA-256 of DCR1E_PREREGISTRATION.md): 628d1b393dff6d074e50484729911ad545dc7313b3f1fe5dee0659b9a1481a1a.