

When Pixels Beat Embeddings: Three Failed Neural Approaches to Symmetry Group Discovery

Author. Jawaun Brown.

Abstract

The Learning the Group paper [2] showed that a simple enumerative procedure — for each candidate rotation θ in a 24-angle grid, score by mean pixel-cosine to the nearest same-class training example — recovers the true Z_8 rotation group from partial-orbit training data at **89.7% recall, 71.3% precision**, with no oracle access. That paper's main limitation, raised in its §6, is that the candidate set must be enumerable. The natural follow-on is to replace enumeration with a *neural* generator that proposes continuous-parameter transformations. We test three natural neural approaches on the same benchmark and find that **all three underperform the pixel-cosine baseline**. We diagnose each failure and report this as an honest null result. The take-away: neural sophistication does not automatically improve symmetry discovery — on controlled rotation data, pixel cosine is essentially the right similarity, and learned features over-generalize in distinctive ways.

1. Introduction

Group-equivariant learning [3, 5] either bakes in the symmetry through architecture, or assumes a known group action. The Learning the Group paper [2] showed the symmetry group can be *recovered from training data* by an enumerative procedure: score each candidate rotation by pixel-cosine self-consistency on training pairs of the same class, keep those above a threshold. On rotated-stroke partial-orbit data this recovers Z_8 cleanly. The paper's own limitations section flagged the enumerable-candidate-set assumption: "the procedure is enumerative... a neural version is needed for non-enumerable groups [6]."

This paper tests three natural neural alternatives:

1. **Direct rotation-angle generator:** a small encoder + latent-conditioned MLP that outputs continuous rotation angles, trained with intra-class consistency + diversity losses.
2. **Encoder-invariance scoring:** a contrastive-trained encoder, then score candidate rotations by mean cosine between $\text{encoder}(x)$ and $\text{encoder}(\text{rotate}(x, \theta))$.
3. **Encoder-based enumerative:** replace pixel cosine in the v1 procedure with cosine in the learned encoder space.

All three **underperform** the pixel-cosine baseline. We document each failure and discuss what it tells us about when neural approaches help group discovery.

2. Method and Setup

We use the rotated-stroke partial-orbit benchmark from [2]: 8 classes, Z_8 rotation, each class shown at only 3 of 8 rotations during training. 8 samples per class $\times$ rotation = 192 training images. Group recovery is measured by recall and precision vs the oracle Z_8 at 7.5° matching tolerance.

2.1 Approach 1: Direct rotation-angle generator

$$G_\varphi(x, z) \rightarrow \theta$$

A small CNN encoder maps the anchor image to a feature; a latent-conditioned MLP combines the feature with $z \in \mathbb{R}^4$ and outputs a scalar rotation angle θ . We rotate the anchor by θ via differentiable `affine_grid + grid_sample` and train with:

- **intra-class**: nearest-same-class pixel-MSE after rotation, averaged over $K = 8$ simultaneous z draws.
- **diversity**: penalize pairs of z draws that produce similar angles, via mean $1 + \cos(\Delta\theta)$.

2.2 Approach 2: Encoder-invariance scoring

A SupCon-style supervised contrastive encoder is trained to pull same-class images together. We define the encoder's *implicit invariance group* as the set of rotations θ such that $e_\phi(\text{rotate}(x, \theta)) \approx e_\phi(x)$ in cosine space. Score each candidate angle by mean cosine over a held-out batch.

2.3 Approach 3: Encoder-based enumerative

Identical to v1 enumerative procedure from [2], but with cosine taken in the SupCon encoder's embedding space instead of raw pixel space.

3. Results

3.1 Approach 1 — mode collapse

Across 128 z draws on 16 anchors, the generator's predicted angles cluster within $\pm 10^\circ$ of identity. After NMS, the single top peak is at 2.5° .

Metric	Value
Top NMS peak	2.5°
Recall vs Z_8	0.125 (only identity recovered)
Precision vs Z_8	1.000 (the one peak is near identity, which is in Z_8)

Diagnostic: the intra-class loss is already low at $\theta = 0$ — the nearest same-class example has some baseline pixel similarity to the anchor without any rotation. The diversity term cannot escape this local minimum without destructive regularization.

3.2 Approach 2 — perceptual smoothness, not data symmetry

The encoder is most invariant to **small rotations near 0°** (perceptual smoothness), not to the Z_8 angles.

Angle	Encoder cosine invariance	In Z_8?
0°	1.000	✓
5°	0.918	✗
355°	0.897	✗
10°	0.806	✗
350°	0.784	✗

15°	0.748	$\times$
20°	0.706	$\times$
345°	0.703	$\times$

Of the top 12 most-invariant angles, only 2 (0°, 320°) are within 7.5° of any Z_8 angle. The encoder has learned to cluster same-class examples into a single embedding point; small rotations preserve that point because they don't change perception much; large rotations move away because perception changes.

Encoder invariance reflects perceptual smoothness, not data symmetry.

3.3 Approach 3 — learned features hurt enumerative discovery

Threshold τ	Kept angles	Recall vs Z_8	Precision vs Z_8
0.5	8	0.375	0.375
0.6	5	0.250	0.400
0.7	3	0.125	0.333
0.8	1	0.125	1.000


For reference, **v1 with pixel cosine** at $\tau = 0.5$: recall **0.897**, precision **0.713**.

Encoder features over-generalize: they treat non-Z_8 rotations like 15°, 30°, 330°, 345° as similar enough to identity that the cosine score is high regardless. Pixel cosine, paradoxically, is sharper because it is literal: rotating by 45° produces a pixel pattern that matches another training example only when 45° is in the true symmetry orbit.


3.4 Headline comparison

Method	Recall vs Z_8	Precision vs Z_8
v1 enumerative with pixel cosine [2]	0.897	0.713
Approach 3 (enumerative + encoder cosine)	0.375	0.375
Approach 2 (encoder invariance scoring)	0.250	0.250
Approach 1 (direct rotation generator)	0.125	1.000

Group recovery on rotated-stroke partial-orbit:
three neural approaches all underperform the pixel-cosine baseline



Approach 2: encoder invariance vs rotation angle.
Peaks at 0° and 180° vicinity; Z_8 angles (red dashes) are NOT preferred.



4. Discussion

The enumerative pixel-cosine procedure works precisely because:

1. It tests *specific candidate rotations* against the training data.
2. The criterion is *exact intra-class match*: $\text{rotate}(x_i, \theta)$ must land near a *literal* training image x_j of the same class.
3. The criterion is *selective*: rotations that don't preserve literal pixel content rarely match anything by chance.

The three neural alternatives fail in distinctive ways:

- **Generators** must output a *parameter*. The loss landscape on this dataset rewards identity-like outputs because same-class examples are already close in pixel space; diversity terms cannot escape the local minimum without overpowering the consistency objective.
- **Encoder invariance** measures perceptual smoothness of the learned representation, not the geometric symmetry of the data. These are different things.

- **Encoder-based enumerative** retains the v1 selection mechanism but loses the v1 similarity sharpness. Learned features blur the geometric precision that pixel cosine exploits.

This does NOT show that neural approaches *cannot* work for symmetry discovery. It shows that two natural neural approaches and one hybrid do not work on *controlled* rotation data where pixel cosine is essentially optimal. The natural follow-on is to test on natural images where pixel cosine is known to fail — cluttered backgrounds, texture variation, lighting, partial occlusion. There, the encoder's ability to ignore irrelevant variation may matter more than its tendency to over-generalize.

5. Limitations

1. Tested only on 16×16 synthetic stroke patterns. Natural-image transfer is untested.
2. Three architectures tried. A conditional flow [van der Ouderaa et al. 6] or a meta-learned equivariance scheme might succeed where these fail.
3. The diversity term in Approach 1 may admit a better hyperparameter that escapes mode collapse; we did not exhaustively sweep.
4. The encoder was a single-architecture supervised-contrastive model; deeper or self-supervised encoders may give sharper invariance.
5. We do not propose a new method that works on this benchmark — only document what doesn't.

6. Reproducibility

```
# Approach 1: direct generator + ensemble
python3 -c "from experiments.neural_group_generator.generator import *; ..."

# Approach 2: encoder + invariance scoring
# Approach 3: encoder-based enumerative discovery
python3 -c "from experiments.neural_group_generator.encoder_invariance import *; ..."
```

See `experiments/neural_group_generator/results/negative_results_2026_06_09.md` for the full result log.

7. References

- [1] **Bennett, M. T.** *How to Create Conscious Machines*. arXiv:2403.00644 (2024).
- [2] **Brown, J.** *Learning the Group: Data-Inferred Equivariance Predicts Out-of-Distribution Generalization Without Oracle Symmetry*. Companion paper (2026). Introduces the v1 enumerative pixel-cosine baseline that this paper compares against.
- [3] **Cohen, T. and Welling, M.** Group Equivariant Convolutional Networks. *ICML* (2016).
- [4] **Khosla, P. et al.** Supervised Contrastive Learning. *NeurIPS* (2020). The SupCon loss we use for the encoder in Approaches 2–3.
- [5] **Kondor, R. and Trivedi, S.** On the Generalization of Equivariance and Convolution in Neural Networks to the Action of Compact Groups. *ICML* (2018).
- [6] **Van der Ouderaa, T. F. A., van der Wilk, M., and Welling, M.** Learning Layer-wise Equivariances Automatically using Gradients. *ICLR* (2024). A neural alternative we did not implement here; the natural Approach 4 for future work.