


Erratum E1: A Perfect Inverted Oracle in the Concern-Gated Retrieval Fixtures

Human director: Jawaun Brown **Producing agent:** Claude Code, directed **Program:** Concern-Gated Retrieval (COGR) **Status:** erratum. Corrects the record of PRs #411, #412, #413, #414, #415. **Date:** 2026-07-24

Summary

After the concern-gated retrieval arc's five papers were merged, a routine information-content check on the concern channel returned an extreme answer: sorting candidates by **ascending** care_anchors and taking the first achieves **hit@1 = 1.000** on every family tested, including the confirmatory pool Wave 1b ran on. care_anchors is a policy-visible field, so a one-line policy outperforms every mechanism the program built.

This erratum records the defect, its root cause, precisely what it does and does not invalidate, the repair, and the one-sort gate that would have caught it at Wave 0. The frozen packages are deliberately left unedited so the historical record and its analysis hashes stay intact and the defect stays visible.



Before and after the prior repair

Found: 2026-07-24, after PRs #411, #412, #413, #414, #415 were merged. **Severity:** program-wide fixture defect. Affects Wave 0, Wave 1a, Wave 1b, MX1. **Frozen packages are NOT edited.** wave0/ and wavelb/ stand as the historical record with the defect documented; their analysis hashes remain valid for what they hashed. The repair lives here, alongside.

1. The defect

Sorting candidates by **ascending** care_anchors and taking the first achieves **hit@1 = 1.000**:

family	bucket	descending hit@1	ascending hit@1
--------	--------	------------------	-----------------

wave0/delayed_commitments	calibration	0.0000	1.0000
wave0/maintenance_fault	calibration	0.0000	1.0000
wave1b/delayed_commitments_v2	calibration	0.0000	1.0000
wave1b/delayed_commitments_v2	confirmatory	–	1.0000 (300/300)
wave1b/maintenance_fault_v2	calibration	0.0000	1.0000

care_anchors is a **policy-visible field on EpisodeContext** – handed to every policy by design. So a one-line policy outperforms every mechanism the program built.

2. Root cause

Wave 0 PREREGISTRATION.md §5 specifies the adversarial prior as: inflate a plausible alarm region to $W_ALARM_INIT = 1.0$, and suppress "*at least one true commitment region*" to $W_COMMIT_INIT = 0.05$. The families implement it as:

```
prior[load_bearing] = W_COMMIT_INIT      # exactly ONE node -- and it IS the answer
```

Concern over candidates therefore takes three values – 0.2 baseline (~10 nodes), 1.0 alarm (~2 nodes), 0.05 suppressed (**exactly 1 node, the answer**). The suppressed value uniquely identifies the target. The prior is adversarial as intended; it is also a perfect inverted oracle, which was not intended.

3. Why every gate missed it

utility, _answer_key. care_anchors is none of those – it is a legitimate policy input, so nothing fired.

learned geometry (learn_graph). The leak is in the hand-authored prior, which G9 never inspected.

(recency / value / priority / salience / wrong-agent). It never tested an **inverted** reading of any signal, including concern itself.

- **G0 IntegrityAudit** forbids policy code from dereferencing role,
- **G9 leakage audit** (label-permutation + randomized-generator) audited the
- **Wave 1a G3 specificity** compared *information-matched* signals

The common shape: every gate asked "can the policy reach something forbidden?" None asked "does something *permitted* already contain the answer?"

4. Validity impact – what stands, what changes

Stands (leak cancels; concern held constant across both arms)

geometry with FROZEN_WRONG concern **identical in both arms**. The leak is present equally on both sides and cancels in the paired difference. "*Learned geometry $\approx$ degree-matched random*" is unaffected.

concern constant.

controls). It scores sets against planted bundles and never touches concern ranking. Unaffected.

- **Wave 1b L1 KILL**. The contrast is LEARNED vs FREQ_MATCHED_RANDOM
- **Wave 1b G2** interventional edge-ablation – a geometry intervention with
- **Wave 1b G9's conclusion about the geometry** – correct for what it audited.
- **MX1 Part B** (verifier split, 109/109 $\rightarrow$ 0, precision 1.0 on 6,920

Re-explained, not invalidated (the leak supplies the missing mechanism)

matched-budget baseline. multiplicative_ppr **upweights** concern, and high concern is precisely the non-answer. The mechanism was not merely using a wrong prior; it was using a *perfect* signal with **inverted sign**. That is why it did so badly, and it is a sharper statement than the original.

is steered actively away from the target. Its 0.545 success rate came from exhausting ~6 of ~13 candidates, not from the ranking.

- **Wave 0's headline** – multiplicative_ppr sitting ~10× *below* the best
- **MX1 Part A.** concern_sequential ranks by multiplicative_ppr, i.e. it

Newly questionable

substrate where a one-line policy scores 1.000. Success rates such as MX1's 0.545 mean "these policies ignored a perfect signal in their own input," not "the task is hard."

oracle ceiling is less remarkable once min-concern *is* the oracle. The specificity slate should have carried an inverted-concern comparator.

weak second signal to be combined with context; it was an oracle used backwards.

- Every **absolute** performance number in the program was measured on a
- **Wave 1a's G3 specificity KILL** – info_matched_recency reproducing the
- The **two-flashlight thesis was never tested as posed**. Concern was not a

5. The repair

prior_repair.repair_wrong_prior suppresses a **set** of k candidates – the load-bearing node plus k-1 non-answer, non-alarm distractors – instead of the answer alone. This honours the preregistration's actual wording ("suppress at least one true commitment region") while removing identifiability. Measured at k = 4, 300 episodes per family:

family	before	after repair
wave0/delayed_commitments	1.0000	0.0633
wave0/maintenance_fault	1.0000	0.1867
wave1b/delayed_commitments_v2	1.0000	0.3767
wave1b/maintenance_fault_v2	1.0000	0.3000

All fall below the 0.8 leak threshold.

Residual, stated honestly. The ideal post-repair value is $1/k = 0.25$; two families land above it (0.3767, 0.3000). The cause is deterministic tie-breaking by node id inside the suppressed set, which leaks a little positional information. A stricter repair would randomise tie-breaking. The current repair is sufficient to clear the gate but is **not** a fully-randomised control.

6. The gate that would have caught it

inverted_signal_audit.audit_signal scores every policy-visible signal s in **both** directions against the answer key and flags a leak if *either* reaches `hit@1 >= 0.8`. One sort per signal per direction; seconds to run.

Standing rule for this program: at fixture-freeze time, audit every policy-visible field in both orderings. Every gate the program had asked whether the policy could reach something *forbidden*. None asked whether something *permitted* already contained the answer.

7. Reproduce

```
uv run --no-sync python -m experiments.concern_gated_retrieval_e2.erratum_e1.verify_erratum
```

Receipt: results/erratum_receipt.json. Local CPU, seconds, no Modal.