

DCR1c: All Six Gates Pass – and the Facet Einstein Actually Deleted Is Absent

Human director: Jawaun Brown **Producing agent:** Claude Code, directed **Program:** Date-Cut Retrodiction – DCR1c **Status:** H1–H6 all **G0** – overall **G0**. DCR2 licensed. **Date:** 2026-07-26

Abstract

DCR1b passed five of six gates and named its own repairs: a polarity and referent test on the privileged-frame matcher, and a third sandboxed extraction pass so consensus is 2-of-3 rather than 2-of-2. DCR1c applies exactly those two and nothing else. Every threshold is carried over as an imported constant.

All six gates pass. Consensus retention rises from 69.3% to **79.7%** (441 propositions). Quote fidelity is 99.8%. Vocabulary residue is 2.32–3.67%. The 1880 deep placebo is silent under **all three** independent extraction passes and under the consensus. And H5 – the gate that failed DCR1b – now returns **12 of 12 genuine** on an individual read.

This is the first clean sweep on material nobody authored for this framework. **DCR2 is licensed.**

Two things temper it, and both are in the results rather than the footnotes.

First, the preregistration was written *after* the run. That is the DR3 slip. What makes it survivable is that no gate was invented, moved or renamed – H1–H6 are definitionally identical to DCR1b's, and run_dcr1c.py imports the thresholds rather than restating them, so they cannot drift. The repairs were specified in DCR1b's published §6 before any of this existed. It is still a confirmatory run by someone who had seen every prior result.

Second, and more interesting: **T1, absolute simultaneity, matches zero propositions at every cut.** The corpus surfaces the privileged aether frame and Lorentz's local time. It does not surface, as an explicit commitment, the thing Einstein actually deleted. That is a finding about the corpus, not a defect, and §6 argues it may be the most informative result here.

1. The two repairs

DCR1b's H5 failed on four T2 hits in three distinct modes, and its paper refused to patch them in place:

mode	example
polarity	"The hypothesis of a stationary ether is shown to be incorrect. "
referent	"...the ether at the earth's surface to be at rest with regard to the earth's surface " – the dragged-ether <i>rival</i>
label	"...in comparison with the fixed system " – a coordinate label

target_v3 adds a polarity veto on targeted refutation markers, a referent veto on rest claims made relative to the earth, and drops fixed system / fixed frame while keeping fixed aether.


The polarity veto deliberately does **not** fire on bare negation, and the reason is a specific sentence in the corpus. Larmor writes: "It has **not** been found possible to construct a system of dynamics which has respect only to the relative positions of moving bodies." That is a negative sentence whose content is exactly the absolute-space commitment. A naive negation veto would delete the strongest kind of evidence the corpus offers.

Vetoos apply to **T2 only**. T1 and T3 produced no false positives under adjudication, and applying an untested veto to a clean facet would be changing something that is not broken.

Validation came before use: v3 was run against DCR1b's adjudicated set – the four hits it must reject and the ten it must keep – and got all fourteen right. That check is a committed regression test, not a claim.

Consensus 2-of-3. A third sandboxed pass was run and consensus now requires agreement in 2 of 3 rather than 2 of 2. Pass 1 remains excluded: its prompt did not forbid reading other repository files and one agent read this repository's own code.

Unchanged: the corpus, the cuts, the extraction prompt, the quorum, every threshold, and residue_v2. DCR1b's modules are not edited.



Repairs


2. Results

gate		
H1 quote fidelity	GO	99.8% vs 90%
H2 vocabulary residue	GO	2.32–3.67% vs an unchanged 5%
H3 deep placebo silent	GO	the 1880 cut matches zero propositions to any facet
H4 target cut not silent	GO	T2 (11) and T3 (1) present at 1904
H5 matcher soundness	GO	12 of 12 genuine on an individual read
H6 robustness across passes	GO	identical verdict under all three passes

Overall GO. DCR2 licensed.

cut	docs	props	residue	facets
1880 deep placebo	3	96	3.67%	none

1897 near placebo	8	227	2.32%	T2
1904 target	15	441	2.37%	T2, T3
1904 no-risk	13	379	2.52%	T2, T3



H6 across three passes. Pass 2, pass 3, pass 4 and the consensus all give: 1880 → nothing, 1897 → T2, 1904 → T2+T3. Given only Maxwell, a model that has read the twentieth century produced 96 consensus commitments and not one touched a privileged frame or local time – under three independent extractions.

That is the circularity COGR Wave 1a died of, tested three times, absent each time.

The third pass did real work. Larmor 1897's *"The spacial framework in absolute rest introduced by Newton is in fact the quiescent underlying æther"* – the single most direct statement of the privileged-frame commitment anywhere in the corpus – is **new at 2-of-3**. DCR1b's 2-of-2 filter dropped it. That is exactly the failure mode DCR1b predicted a third vote would fix, caught in the act.

3. Freeze status, and why this is the weakest of the three

DCR1C_PREREGISTRATION.md was written **after** run_dcr1c.py was executed. That is the DR3 slip and I am not dressing it up.

What is defensible:

DCR1b's, and the runner **imports** QUOTE_FIDELITY_GATE and RESIDUE_RATE_GATE from run_dcr1.py rather than restating them, so they cannot drift.

existed. DCR1c implements those two and nothing else.

before being run on new data.

- No gate was invented, moved or renamed. H1–H6 are definitionally identical to
- The repairs were specified in DCR1b's published paper §6, before any of this
- target_v3 was validated against the previous paper's adjudicated failures

What that does not buy: this is a confirmatory run of a hypothesis I already suspected, by someone who had seen every prior result. The one thing it cannot be is threshold-fitted.

4. Adjudication

All twelve hits at the target cut were read. All twelve state their facet.

Two are **conditional** in form – "If the ether is at rest and the apparatus moves..., the directions and distances traversed by the rays are altered" – and I counted them genuine because the antecedent states the theory's commitment and the consequent draws a physical consequence. They are not *reductios*; the *reductio* case ("then Lorentz's own theory also fails") was correctly vetoed. Both judgments are flagged in `results/dcr1c_facet_adjudication.json` so the call is visible rather than buried.

A fragility no gate measures. T3 still rests on a *single* proposition – Lorentz's "The transformed time variable t' may be called the local time." The target cut clears its quorum of two by a margin of one. If that one proposition had been missed, H4 would have failed. Nothing in DCR1b or DCR1c repaired this, and no gate reports it. A successor should either widen T3 or add a margin gate.

5. What is now established

Across DCR1, DCR1b and DCR1c, on a corpus of 15 public-domain documents from Maxwell 1865 to Lorentz 1904:

contamination was found and stripped structurally before extraction.

translated in 1913 – was chased to the French originals and cleared.

corpus's own words.

- The corpus is clean and checksummed. Wikisource's Portal:Relativity header
- Its one serious provenance risk – every sentinel term riding on two documents
- Quote fidelity is essentially perfect across every pass.
- Vocabulary residue, measured properly, is 2.3–3.7%: the extractor works in the
- **The deep placebo is silent under three independent extractions.**
- Every facet hit at the target cut survives an individual read.

6. The absent facet, which may be the real finding

T1 – absolute simultaneity – matches **zero** propositions at every cut, under both v2 and v3, across all three passes. It is not that the matcher is too strict: v3's T1 was validated to fire on four genuine phrasings including "There is an absolute time common to all systems" and "The time is the same for a stationary observer as for an observer carried along in uniform motion."

The corpus does not contain those sentences. It contains the privileged aether frame, stated many times and many ways. It contains local time, named by Lorentz as a transformed variable. It contains Poincaré stating the principle of relativity in September 1904 – **and keeping the ether.**

What it does not contain is anyone writing down "simultaneity is absolute" as a commitment that might be given up. That is precisely what one would expect of a presupposition so deep that nobody thought to state it, and it is why the deletion was available to be made and was not made for another nine months.

This has a sharp consequence for DCR2, and it is not a comfortable one. If the load-bearing deletion is a proposition **no document states**, then nomination over an extracted proposition set cannot rank it – it is not in the candidate set at any rank. DCR2 as currently conceived can measure whether the nominators find the *privileged frame*, which the corpus does surface.

It cannot measure whether they would have found *absolute simultaneity*.

Two honest readings, and I do not think this paper can decide between them:

instrument – one that infers commitments from what the reasoning *requires* rather than from what the text *says*. The framework's five-slot schema already distinguishes asserted from presupposed; the extractor's kind: presupposed field exists and is evidently not reaching this deep.

nobody states, then deletion-repair nomination over an extracted set is structurally unable to find them, in the same way DR2 proved the two-nominator claim unreachable under its original cost definition. That would be a theorem-shaped result and it deserves to be pursued as one.

1. **A limit of extraction.** Unstated presuppositions need a different
2. **A limit of the framework.** If the deletions that matter are the ones

DR2's precedent argues for taking possibility 2 seriously rather than assuming possibility 1. The way to separate them is to ask whether *any* extraction can surface a presupposition the corpus never states – a question with a clean experimental shape, and the one I would run before DCR2 rather than after.

7. Scope

DCR2 is licensed and can be run. Nothing here says the nominators work on real material; that remains its question. Nothing here bears on vocabulary extension, the standing ceiling of the whole framework. And §6 argues that DCR2, run as specified, would answer a narrower question than the programme was aiming at.

Appendix: reproduction

```
uv run --no-sync python -m experiments.date_cut_retrodiction.run_dcr1c
```

Local CPU, seconds; it rebuilds the 2-of-3 consensus before scoring. `residue.py`, `target.py`, `target_v2.py`, `residue_v2.py` and `consensus.py` are **not** edited – DCR1's and DCR1b's published numbers remain reproducible, and `target_v3.compare_v2_v3` quantifies what the vetoes change.

Figures: `papers/date_cut_retrodiction_dcr1c/figures/build_figures.py`.