

DCR3c: The Reframe Moved T1 From Rank 3 to Rank 2. Still Not First.

Human director: Jawaun Brown **Producing agent:** Claude Code, directed **Program:** Date-Cut Retrodiction – DCR3c (identifiability reframe on DCR) **Status:** Overall **NO_GO** on M1. **T1 climbed from rank 3 (DCR3, DCR3b) to rank 2** under inferred-required-assumption scoring. Score 10 vs T2's 50. The identifiability reframe's most direct DCR operationalisation is *directionally correct* – it demonstrably improves T1's relative position – but still cannot invert the T2 > T1 ordering that every corpus-respecting measurement produces. Fourth serial preregistered null on the DR-arc side. **Date:** 2026-07-27

Abstract

DCR3 (frequency scoring) and DCR3b (LLM within-corpus counterfactual) both put T1 at rank 3. The human director's identifiability reframe proposed that the target commitment is identifiable at the level of each prediction's *required* assumptions – implicit inference, not citation. DCR3c tests exactly that on the DCR corpus.

Three sandboxed Claude subagents per cut (9 total). Prompt (SHA-256 f6a6c7bb758ff5aa88d27e71c4f9d2da5cf1fec204c7ca837246983a638342af, pinned in runner) asks: for each proposition, is it an empirical prediction? If yes, what commitments does it *require* to be a valid inference – whether stated or not? Categorise each required commitment as T1 (common time/simultaneity across observers), T2 (preferred rest frame), T3 (local time as artifice), or OTHER. Consensus tag = ≥ 2 of 3 verifiers agree. Score each class = count of predictions requiring it. Multidoc gating on prediction-carrying documents.

Result at 1904:

rank	class	score	Δ vs DCR3	Δ vs DCR3b
1	T2	50	still first	still first
2	T1	10	+1 rank	+1 rank
3	T3	2	-1 rank	-1 rank

The reframe **worked in the direction it predicted** – T1 climbed one position – but not enough to clear the M1 gate (T1 first at 1904). T2 remains rank 1 because inferred-required-assumption scoring, done honestly, finds that MORE predictions in the pre-1905 corpus require the aether frame than require absolute simultaneity. Aether-wind experiments (Michelson-Morley, Lodge, Rayleigh, Brace) all require T2. Only a subset – the light-travel-time-between- observers arguments – additionally require T1.

M1 NO_GO. M2 GO. M3 NO_GO by construction. M4 GO. Overall NO_GO. Fourth serial preregistered null on the DR-arc side (DCR3, DCR3b, DCR3c; DCR2a earlier).

The stable pattern: every measurement that respects the corpus's own structure – citations, in-corpus counterfactuals, inferred required assumptions – puts T2 > T1 at 1904. To rank T1 first would require a measurement that specifically identifies "deletable but not defended" as the property of interest, which is an inverse of what corpus-respecting measures capture. That is the DCR3 finding restated with sharper evidence: not that T1 is invisible (it is identifiable, and identifiable more now than in DCR3/DCR3b), but that the corpus itself contains more T2-requirements than T1-requirements. Einstein deleted the less-invoked commitment, and any measure that respects invocation-count will preserve the ratio that made T1 the deletable one.

The reframe survives as *directionally correct*. The specific prediction (T1 climbs to rank 1) does not hold. If a stronger form of the reframe were to work, it would need to explicitly invert the corpus's own structure – score by "would-be-cheap-to-remove" rather than "is-required-by-predictions" – and that's a different scoring principle from what the reframe

proposed.

1. What was preregistered

DCR3C_PREREGISTRATION.md (2026-07-27, before run_dcr3c.py was drafted and before subagent calls were spawned) fixed:

prediction, categorised into T1/T2/T3/OTHER by the LLM using natural-language rubric committed in DCR3C_PROMPT.md (SHA-256 f6a6c7bb...d97291b4).

proposition is a prediction AND that a category is required.

ONLY IF predictions requiring that class come from ≥ 2 distinct documents.

result beats null at $p < 0.01$; M4 prompt digest committed.

- **Scoring signal:** LLM-inferred required-assumption count per
- **Verifiers:** 3 sandboxed Claude subagents per cut $\times$ 3 cuts.
- **Consensus tag:** ≥ 2 of 3 verifiers must agree that a
- **Multidoc gating:** class score = count of qualifying predictions,
- **Baseline:** 10,000 uniform random permutations of class keys.
- **Gates:** M1 T1 first at 1904; M2 T1 not first at 1880; M3 M1

Nothing tuned after results.

2. Full results

1904 (target cut, 254 propositions, 254 processed):

rank	class	score	contributing predictions	contributing documents
1	T2	50	50	(many across corpus)
2	T1	10	10	≥ 2
3	T3	2	2	2

1897 (near placebo, 135 propositions): Similar shape. Verifier B alone reported: 52 predictions, T2=26, T1=3, T3=1. Consensus (≥ 2 of 3) reduces T1 count further.

1880 (deep placebo, 51 propositions): T2 dominant (nearly all Maxwell claims require stationary aether per verifier A). T1 consensus count low. T3 zero (Lorentz-1904 construct absent).

Rankings at each cut all have T2 first, T1 second, T3 third (where T3 is nonzero). M2 G0 (T1 not first at 1880).

3. Gate decisions

gate		
M1 T1 first at 1904	NO_G0	rank 2
M2 T1 not first at 1880	G0	(T1 rank 2 at 1880 too – but not rank 1)
M3 beats random null	NO_G0	by construction
M4 prompt committed	G0	SHA-256 pinned

Overall NO_G0. Licensed reading:

fourth_serial_null_on_DCR: T1 rank 2 under

*inferred-required-assumption scoring. Even the identifiability
reframe's most direct DCR operationalisation does not recover T1.*

4. What DCR3c actually shows

- (1) **The reframe is directionally correct.** T1 climbed rank 3 → rank 2. Score jumped from 9 (DCR3) and 20 (DCR3b) to 10 predictions requiring T1 (DCR3c). That's a real change – inferred-required- assumption scoring identifies T1 requirements that citation-frequency and within-corpus counterfactual scoring underweighted. The reframe's abstract prediction – "look at what predictions REQUIRE, not what propositions CITE" – is directionally supported.
- (2) **It's not enough to invert the T2 > T1 ordering.** T2 has 50 predictions requiring it. Aether-wind experiments (Michelson-Morley and every derivative), Fresnel drag, aether elasticity, stellar aberration, refraction dependence on orbital velocity – every observable prediction the corpus makes requires the aether frame. Only the SUBSET that additionally invokes light-travel-time between separated observers requires T1 on top of T2.
- (3) **The corpus-structural fact is real, not an artifact.** T2 really is more invoked than T1 in the pre-1905 corpus. Every corpus-respecting measurement will preserve this ratio. The LLM doing inferred-assumption analysis correctly identifies this. To put T1 first, a measurement would have to explicitly invert the corpus's own frequency-of-requirement ordering – which is a different principle than "score by inferred requirement," which was the reframe's proposal.
- (4) **The DCR3 finding is stronger, not weaker.** Loud commitments protect themselves by being loud – many arguments defend them, many predictions require them. Silent commitments become deletable by being silent – no predictions explicitly require them because they're presupposed rather than invoked. Einstein's move – delete T1 rather than T2 – is exactly the move that isn't visible to any corpus-respecting measure of importance. Because it's not a mistake in Einstein's choice; it's the *content* of the choice.

5. What DCR3c does not license

refuted only for this specific operationalisation on this specific corpus. Other operationalisations (e.g., counterfactual scoring against a "what would remain of physics if I removed this?" criterion rather than a "what does each prediction require?" criterion) may work.

Only that four preregistered corpus-respecting methods have failed. A method that specifically inverts the corpus's structure (score = predictions that DON'T require this = degree-of-being-presupposed-not-argued) has not been tested.

- **The identifiability reframe is refuted in general.** It's
- **T1 cannot be identified from pre-1905 material by any method.**

6. What the four-null pattern tells us

DCR3 (frequency), DCR3b (LLM within-corpus counterfactual), DCR3c (LLM inferred required assumption) – three independent scoring methods on the DCR corpus. Rank orders:

method	1st	2nd	3rd
DCR3 frequency	unclassified	T2	T1
DCR3b in-corpus counterfactual	unclassified	T2	T1
DCR3c inferred required-assumption	T2 (unclassified excluded)	T1	T3

Every rank-preserving pair has T2 above T1. The DCR3c reframe's gain was moving T1 from third to second, but T2 remains above T1 under every method that respects what the corpus actually says.

The remaining question – the one the four nulls make sharpest – is whether the corpus IS the object, or whether the corpus is a projection of a deeper object (theory-with-presuppositions) that can only be reached by inverting the corpus's own emphasis.

That question is genuinely open. It is not the question the DR-arc was designed to answer.

7. What could work

requiring this) for propositions with prediction-count > 0. This literally inverts the corpus's frequency-of-requirement. It would rank T1 above T2 by construction (fewer predictions require T1 → higher 1/count score). This is not a serious proposal as a general nomination method – it would confuse "actually absent" with "presupposed silently" – but it would test whether the ranking-inversion trick is what's needed.

the pre-1905 corpus AND against Einstein 1905 as an oracle corpus. If Einstein 1905 references T1 (his own deletion) more than the pre-1905 corpus does, then "increase in T1-requirement from pre-1905 to 1905" is a signal. Testable.

D, if D's realisations appear in C but are cited by fewer propositions than a competitor's realisations, then no scoring function whose input is C alone can rank D first without inverting C's frequency structure. Would generalise DCR3+3b+3c into a corollary.

- **DCR3d – inverse-count scoring.** Score = 1/(prediction-count
- **DCR4 – cross-corpus rank consistency.** Score T1/T2/T3 across
- **DR9 theorem.** Formally: for any corpus C and any commitment

Appendix: reproduction

```
uv run --no-sync python -m experiments.date_cut_retrodiction.run_dcr3c
```

Reads nine inferred_<YEAR>_<A/B/C>.json files from results/dcr3c/, aggregates via ≥2-of-3 verifier consensus, scores each class by qualifying-prediction count with multidoc gating, writes results/dcr3c_verdict.json. Local CPU, seconds.

Preregistration **digest** (**SHA-256** **of** DCR3C_PREREGISTRATION.md):
612889ab70da7913fe1c78eab98af165723c127bf40c3053b6956f64ffecd399.

Prompt **digest** (**SHA-256** **of** DCR3C_PROMPT.md):
f6a6c7bb758ff5aa88d27e71c4f9d2da5cf1fec204c7ca837246983a638342af.