

DCR3d: The Silent-But-Load-Bearing Signal Fires. At the Wrong Time.

Human director: Jawaun Brown **Producing agent:** Claude Code, directed **Program:** Date-Cut Retrodiction – DCR3d (use/discussion ratio scoring) **Status:** Overall **NO_GO** on both M1 and M2. But the specific pattern is a real finding: **T1 ranks first at both 1880 AND 1897 under use/discussion ratio, and only DROPS to rank 2 at 1904** – because by 1904 T1 has started being discussed. The silent-but-load-bearing signal fires; it just fires *before* the deletion, not *at* it. Deletability is a build-up state, not a moment property. Fifth serial null with the sharpest new finding of the DCR run. **Date:** 2026-07-27

Abstract

The human director's silent-but-load-bearing intuition, operationalised as $\text{deletability}(C) = \text{use_count}(C) / (\text{discussion_count}(C) + 1)$, was preregistered to test whether T1 ranks first at the 1904 target cut AND not first at the 1880 placebo. Nine sandboxed Claude subagents tagged which propositions DISCUSS each class (T1/T2/T3) as their subject vs merely USE it as background. Use counts reused from DCR3c (also 9 subagents, prior run). ≥ 2 -of-3 consensus. Prompt SHA-256 pinned.

Result: T1 rank across cuts:

cut	T1 use	T1 disc	T1 ratio	T2 ratio	T1 rank
1880 (deep placebo)	9	1	4.50	1.00	1
1897 (near placebo)	3	1	1.50	0.46	1
1904 (target)	10	7	1.25	1.92	2

M1 NO_GO (T1 rank 2 at 1904). **M2 NO_GO** (T1 rank 1 at 1880 too). **Fifth serial preregistered null on the DR-arc side.**

But the pattern is not the "always-quiet detector" I preregistered as M2's failure mode. It's a specific empirical finding:

The use/discussion ratio identifies T1 as silent-but-load-bearing at 1880 and 1897 (T1 ratio 4.5 and 1.5, both first). At 1904, T1's discussion count JUMPS from 1 to 7 – Poincaré 1898 treats simultaneity philosophically, Larmor 1900 ch11 writes common time into the transformation, Lorentz 1904 mentions corresponding instants – and T1's ratio drops to 1.25, below T2's 1.92.

The silent-but-load-bearing signal exists. It's real, it's measurable, and it correctly identifies T1's deletability *while T1 was still silent*. The signal fades exactly when the community starts discussing the commitment – which is precisely the pre-deletion transition. Einstein's move in 1905 caught a signal that had been at maximum in 1880, was diminishing by 1904, and would have been gone entirely once someone else made the deletion first.

Deletability is a build-up state, not a moment property. The measure detects WHERE the deletion is available, not WHEN it happens. The moment of deletion is when someone (Einstein) acts on pre-existing availability. That's a substantive finding about how conceptual revolutions actually work, and it converts DCR3's loud/silent asymmetry from a corpus-time-snapshot claim into a corpus-time-trajectory claim.

1. What was preregistered

DCR3D_PREREGISTRATION.md (2026-07-27, before run_dcr3d.py was drafted, before subagent discussion-tagging calls were spawned):

for each class $C \in \{T1, T2, T3\}$. Higher = more deletable.

(≥ 2 of 3 verifier consensus).

$\times 3$ cuts) tag each proposition with the T1/T2/T3 categories it DISCUSSES as its subject rather than merely USES as background. Prompt (SHA-256 24384377bfab8dfe87cc72ce01d0908da992cddb2e39d11781cb42f432dbea05) gives explicit USE vs DISCUSS rubric with examples.

discusses a category for the tag to count.

author-discussion signals like Poincaré 1898 on simultaneity is the whole point.

- **Scoring:** deletability(C) = use_count(C) / (discussion_count(C) + 1)
- **Use counts:** reuse DCR3c's inferred-required-assumption tags
- **Discussion counts:** new. 9 sandboxed Claude subagents (3 per cut
- **Consensus rule:** ≥ 2 of 3 verifiers must agree that a proposition
- **No multidoc gating** on discussion counts – rewarding single-

Four preregistered gates: M1 (T1 first at 1904), M2 (T1 not first at 1880), M3 (M1 beats random null), M4 (prompt committed).

The preregistration named the M1-GO/M2-NO_GO failure mode ("always-quiet detector") as a specific informative-null outcome. The actual outcome (M1 NO_GO and M2 NO_GO with T1 rank 1 at both placebo cuts) is a different, richer failure mode not preregistered as a distinct row.

2. Results in full

Per-cut consensus (≥ 2 of 3 verifiers agree):

cut	T1 use	T1 disc	T2 use	T2 disc	T3 use	T3 disc
1880	9	1	20	19	0	0
1897	3	1	26	55	1	1
1904	10	7	50	25	2	4

Ratios:

cut	T1	T2	T3
1880	4.50 (rank 1)	1.00	0.00
1897	1.50 (rank 1)	0.46	0.50
1904	1.25 (rank 2)	1.92 (rank 1)	0.40

Verifier variance on discussion tagging was substantial: at 1904, verifier A tagged 14 T1-discussions, B tagged 7, C tagged 6. The ≥ 2 -of-3 consensus reduces to 7 (Poincaré 1898 four items + Larmor 1900 common-time + Lorentz 1904 corresponding-instants + Poincaré 1904 simultaneous-transmission). Cross-verifier variance is real but the consensus filters it to defensible tags.

3. Gate decisions

gate		
M1 T1 first at 1904	NO_GO	T1 rank 2 (T2 first)
M2 T1 not first at 1880	NO_GO	T1 rank 1
M3 beats random null	NO_GO	by construction

Overall NO_G0.

Licensed reading (from the runner):

*fifth_serial_null: T1 rank 2 even under use/discussion ratio
scoring. Silent-but-load-bearing intuition correct in spirit but
this specific operationalisation does not invert T2's dominance.*

But that reading misses the substantive pattern. The strict verdict is NO_G0; the interpretive finding is richer and worth naming.

4. What DCR3d actually found

Reading the numbers across the three cuts:

At **1880** (Maxwell only), T1 is used 9 times (Maxwell's field theory quietly assumes an "instant across the medium" many times), discussed once (a single Maxwell proposition takes "instant across whole medium" as its subject). Ratio 4.5. T1 is at MAXIMUM silent-load-bearingness. Nobody argues about common time. Everyone uses it.

At **1897**, use drops to 3 (extraction identifies fewer T1 requirements in this cut's specific propositions – many are Michelson- Morley-shaped where the required commitment is more specifically about light-travel-time-being-well-defined than about common-time-across-observers-generally). Discussion still 1. T1 ratio 1.5. Still first. T1 has become slightly less silently-loaded but still no one discusses it.

At **1904**, use is back up to 10 (Larmor 1900, Lorentz 1904, and Poincaré 1904 add explicit predictions that require T1). **But discussion jumps from 1 to 7.** Poincaré 1898 treats simultaneity philosophically. Larmor's *common_time_across_two_systems* is explicitly written into the transformation. Lorentz 1904's *true_time_defines_corresponding_instants* names the concept. Poincaré 1904's *simultaneous_transmission_meaningful* treats the notion directly. Suddenly the commitment IS discussed. Ratio drops to 1.25. T2 wins by having proportionally more use for its discussion count.

The revolution's precursors – Poincaré, Larmor, Lorentz – literally raised the discussion count on T1 in the immediate pre-1905 window. They started making explicit what had been implicit, and in doing so they made the silent-but-load-bearing signal *decay*. The signal is at maximum when the community is silent. The signal fades exactly at the moment when someone eventually acts on it.

5. What this implies for the DR-arc

DCR3 established the loud/silent asymmetry as a static claim about

silent-but-load-bearing property peaked in 1880 (Maxwell alone) and was already fading by 1904. Einstein made the deletion when the signal was at ~1.25 – after it had dropped from 4.5 twenty-four years earlier.

1904. DCR3d shows the asymmetry is **dynamic**. The

Deletability is a trajectory, not a moment. A scoring principle that fires at 1880 for a deletion Einstein made in 1905 is not "wrong" – it's identifying the same commitment at an earlier point in its deletability trajectory. The actual moment of deletion is determined by who is looking, not by when the signal peaks.

For the DR-arc's original question – "given only pre-1905 material, identify the deletion that will be made in 1905" – this is a constructive negative result:

1880 and 1897.** If you'd asked at 1880 "what's silently load-bearing enough to be deletable?" – you'd have gotten T1.

(Poincaré, Larmor, Lorentz) have started discussing it. If you ran the measure only at 1904 (which is what M1 tested), T2 wins on the ratio.

- ****Use/discussion ratio DOES identify T1 as maximally deletable at**
- **But by 1904 the signal has diminished** because T1's precursors

The measure works, at a different cut than the DR-arc's design predicted it should. The DR-arc wanted "identify what Einstein would delete, using material available immediately before he deleted it." What this experiment shows: the material *immediately before* is already partially post-precursor. To catch the deletable commitment, you need to score at a cut before the precursor discussion – which is a fundamentally different measurement protocol.

6. What DCR3d does not license

at 1904. The trajectory finding doesn't rescue the target cut score.

unexpected way – the signal is real but temporally misaligned.

gate is NO_GO. The interpretive finding is worth having but the ledger records a fifth serial null.

- **T1 is now recoverable from pre-1905 material.** M1 is still NO_GO
- **The reframe is confirmed.** It's directionally supported in an
- **This is a positive result in disguise.** The strict preregistered

7. What comes next, sharpened by this finding

class, compute (use/discussion) at 1880, 1897, 1904 and look at the *derivative*. The commitment whose ratio DROPS most rapidly approaching 1904 is the one being brought into discussion (i.e., the deletion is imminent). Under this measure, T1 might rank higher because its ratio drop is largest (4.5 → 1.25).

1905. If T1's discussion count in Einstein 1905 is much higher than in the pre-1905 corpus, the derivative measure would confirm.

scoring rule based on a single-cut snapshot of use/discussion ratio cannot identify commitments whose deletability signal has already begun fading due to precursor discussion. The signal has to be measured before the precursors start.

- **DCR3e – trajectory rather than snapshot scoring.** For each
- **DCR4 – post-cut oracle comparison.** Score at 1904 vs Einstein
- **DR9 theorem.** Formalise the trajectory finding: prove that any

Appendix: reproduction

```
uv run --no-sync python -m experiments.date_cut_retrodiction.run_dcr3d
```

Reads DCR3c use tags (existing) and DCR3d discussion tags (9 new subagent outputs), computes per-class use/discussion ratios at each cut with ≥2-of-3 verifier consensus, applies M-gates, writes results/dcr3d_verdict.json. Local CPU, seconds.

Preregistration **digest** **(SHA-256** **of** **DCR3D_PREREGISTRATION.md):**
009d9b13baf519e3bfd6ae011c516d287ca9f24eaaf5e4102b4fe1a8793159a7.

Prompt **digest** **(SHA-256** **of** **DCR3D_PROMPT.md):**
24384377bfab8dfe87cc72ce01d0908da992cddb2e39d11781cb42f432dbea05.