

Translation Augmentation Produces Toroidal Codes and Larger-Arena Generalization in Path-Integration RNNs

Jawaun Brown

Research-Derived Experiments · empirical note and negative mediation result

Abstract

This note tests a pre-registered strong claim about path-integration RNNs: **weakness**, a held-out linear equivariance score for wrapped translations, should govern whether a learned population code forms a toroidal manifold, and toroidal topology should mediate larger-arena generalization. The result is a useful negative mediation result. In a 320-network Modal H100 sweep, full translation augmentation reliably induces toroidal codes (G1 pass; torus match 0.734), preserves decoding when the arena doubles (OOD 0.949 versus 0.484 for no augmentation), and produces the causal condition contrast specified before the sweep (G6 pass). Weakness also tracks spectral concentration (G5 pass; $\rho = +0.635$). But the central triangle fails: weakness only weakly predicts toroidal score (G2 fail; $\rho = +0.197$), does not beat final loss on larger-arena OOD by the pre-registered 2x margin (G3 fail; weakness $\rho = +0.617$, loss $\rho = +0.652$), and topology does not mediate the weakness-OOD relation (G4 fail). The surviving claim is therefore bounded: in this artificial RNN harness, translation-structured training causally produces toroidal, OOD-generalizing codes, while weakness measures one spectral aspect of the learned translation structure rather than serving as the governing scalar of torus formation.

1. Framing

Biological grid-cell population activity has been shown, using topological data analysis, to lie on a toroidal manifold, and path-integration RNNs can learn grid-like spatial codes. This experiment asks a bridge question: when toroidal codes and larger-arena generalization appear in a learned path integrator, are they governed by the same substrate-general weakness quantity that predicted OOD behavior in the symbolic and vision experiments, or are they more directly caused by the translation-structured training condition itself?

The strong version predicted a triangle: high weakness, low Fourier participation ratio (few aligned frequency modes), clean toroidal topology, and larger-arena OOD generalization should be different measurements of one event. This sweep was designed to test that strong triangle and let it fail. The published posture should therefore be an empirical note, not a broad confirmation paper.

2. Pre-Registered Test

The gates were specified in the project repository before the Modal sweep (`papers/grid\_cell\_weakness/preregistration.md`, frozen 2026-06-28). The registration was project-level rather than an external registry filing. It asked whether a single translation-compatibility scalar predicts both population topology and OOD accuracy, and whether topology carries the weakness-to-OOD relation.

Gate	Pre-registered criterion	Modal result	Verdict
G1 manifold	full translation torus match ≥ 0.60	0.734	pass
G2 weakness-topology	$\rho \geq 0.5$ and $\geq 2\times$ best baseline	$\rho +0.197$; loss $+0.431$	fail
G3 weakness-OOD	$\rho \geq 0.5$ and $\geq 2\times$ best baseline	$\rho +0.617$; loss $+0.652$	fail
G4 mediation	partial ρ drops by $\geq 50\%$	partial $\rho +0.623$; no drop	fail
G5 spectrum	$\rho(\text{weakness, -Fourier PR}) \geq 0.5$	$+0.635$	pass
G6 causal	full translation beats none/random shift	topology and OOD both lift	pass
wrong-group null	wrong-group predictor null or negative	tie-aware ρ 0.000	pass

Table 1. Locked gates and outcomes. The decisive failure is not auxiliary: G2-G4 reject the registered mediation triangle.




Figure 1. Gate verdicts from the pre-registered Modal sweep. The causal and spectral legs pass; the topology-governing and mediation claims fail.

3. Methods

We trained 320 velocity-driven path-integration RNNs on Modal H100 workers: five training conditions, two recurrent architectures, and 32 seeds. Each network predicts a place-cell code from velocity sequences. Conditions vary the augmentation regime: full translations, partial translations, random shifts, no augmentation, and a wrong-group control. For each trained code we measure weakness, persistent-homology toroidal score, Fourier participation ratio, and decoding accuracy in held-out arenas of scale 1.0, 1.25, 1.5, and 2.0. The largest arena is the pre-registered OOD metric.

Weakness. For a binned hidden-state population code $H(x)$ and a held-out wrapped translation τ , fit one least-squares linear operator A_τ on half the spatial bins and evaluate it on held-out bins: $W = E_\tau R^2[H(x + \tau), A_\tau H(x)]$. Translations are wrapped modulo the 16×16 bin grid, one operator is fit per translation, R^2 is computed over hidden units and held-out positions, and the reported weakness is the mean clipped R^2 across the registered shifts. The metric is computed on hidden activity, not on the decoded position or place-cell target.

Topology. Hidden states are sampled from fresh trajectories, averaged into a 16×16 spatial grid, and treated as a point cloud in hidden-state space. Empty spatial bins, when present, are filled with the global mean so the grid is complete; the Modal cells report coverage. Persistent homology is computed with a Vietoris-Rips complex (Gudhi backend, max simplex dimension 3, edge cap at the 45th percentile of pairwise distances), yielding H_0 , H_1 , and H_2 persistence intervals. A torus should have Betti signature (1,2,1): one component, two loops, and one void. The continuous toroidal score combines the second H_1 bar above the third-bar noise floor with the strongest H_2 void; ``betti_match_torus`` requires two estimated H_1 loops and a nontrivial H_2 bar.

4. Results




Figure 2. Condition means. Full translation is the only condition that reliably forms a torus and preserves OOD decoding at arena scale 2.0.




Figure 3. Larger-arena OOD curves. Full translation remains stable as the arena doubles; controls degrade sharply.

Condition	n	weakness	toroidal	torus match	OOD @2.0
full translation	64	0.768	0.357	0.734	0.949
partial translation	64	0.416	0.007	0.000	0.732
random shift	64	0.400	0.000	0.000	0.615
none	64	0.446	0.000	0.000	0.484
wrong group	64	0.048	0.009	0.000	0.489

Table 2. Condition means. Full translation is the positive intervention; random shift and wrong group are controls. Weakness is not monotone with toroidal score outside the full-translation condition.

The positive result is clean at the condition level. Full translation augmentation is the only condition that reliably produces a torus, and it preserves decoding accuracy as the arena doubles. Partial translation and random shifts lift OOD relative to no augmentation, but they do not

produce the torus match. The wrong-group control has near-zero weakness under the registered translation metric and does not lift topology or OOD.

The negative result is equally important. Across individual networks, weakness does not explain toroidal score strongly enough to pass the registered gate, does not outperform final loss on OOD by the required margin, and leaves the weakness-OOD association intact after controlling for topology. These are not wording problems; they falsify the strongest version of the theory tested here.

5. Interpretation

The result should not be polished into a full confirmation. It is better than that: it marks the boundary of the weakness program on a task where topology is genuinely load-bearing. Translation augmentation is the experimentally manipulated variable that produces toroidal topology and larger-arena generalization in this harness. Weakness is still useful as a spectral and OOD-associated measurement, but it is not the scalar that governs torus formation here.

A plausible reading is that the augmentation condition itself is the strongest causal variable: it directly supplies the group structure needed for both topology and OOD. Weakness measures part of that event, especially its spectral footprint, but it is not sufficiently specific to explain topology variance after training loss and condition effects are present. The failed gates are not auxiliary failures; they reject the strongest claim. The surviving claim is not that weakness governs toroidal topology, but that translation-structured training produces toroidal, OOD-generalizing codes while weakness tracks one spectral aspect of learned translation structure.

6. Scope and Limitations

This remains an artificial RNN path-integration task, not a biological replication. The paper should not claim to explain biological grid cells. Its claim is narrower: a learned path-integration model, inspired by the grid-cell/topology literature, responds to a translation-augmentation intervention by forming toroidal hidden-state structure and generalizing to larger arenas. The topology metric uses binned hidden activity and Vietoris-Rips persistent homology; both can be sensitive to sampling density, binning, edge thresholds, and seeds. The wrong-group null in the first Modal summary was miscomputed by a rank helper that did not average ties; the corrected tie-aware Spearman is reported here and in the result report. The raw cell measurements did not change.

Future tests should separate condition effects from within-condition variation by sampling a larger family of already-toroidal codes, perturbing trained codes rather than only training regimes, and checking topology robustness under multiple bin counts, point sampling strategies, and persistent-homology backends. A biological extension would be a new confirmatory study on public grid-cell recordings, not an interpretation of the RNN result itself.

References

- Gardner, R. J. et al. Toroidal topology of population activity in grid cells. *Nature* 602, 123-128 (2022).
- McNaughton, B. L., Battaglia, F. P., Jensen, O., Moser, E. I., Moser, M.-B. Path integration and the neural basis of the cognitive map. *Nature Reviews Neuroscience* 7, 663-678 (2006).
- Fuhs, M. C., Touretzky, D. S. A spin glass model of path integration in rat medial entorhinal cortex. *Journal of Neuroscience* 26, 4266-4276 (2006).
- Burak, Y., Fiete, I. R. Accurate path integration in continuous attractor network models of grid cells. *PLoS Computational Biology* 5, e1000291 (2009).
- Cueva, C. J., Wei, X.-X. Emergence of grid-like representations by training recurrent neural networks to perform spatial localization. *arXiv:1803.07770* (2018).

- Banino, A. et al. Vector-based navigation using grid-like representations in artificial agents. *Nature* 557, 429-433 (2018).
- Sorscher, B., Mel, G. C., Ganguli, S., Ocko, S. A. A unified theory for the origin of grid cells through the lens of pattern formation. *NeurIPS* (2019).
- Xu, D., Gao, R., Zhang, W.-H., Wei, X.-X., Wu, Y. N. Emergence of grid-like representations by training recurrent networks with conformal normalization. *arXiv:2310.19192* (2023).
- Cohen, T., Welling, M. Group Equivariant Convolutional Networks. *ICML* (2016).
- Gruver, N., Finzi, M., Goldblum, M., Wilson, A. G. The Lie Derivative for Measuring Learned Equivariance. *ICLR* (2023).
- Xu, M., Song, F., Si, B., Qin, S. The Principle of Isomorphism: A Theory of Population Activity in Grid Cells and Beyond. *arXiv:2510.02853* (2025).
- Bennett, M. T. How to Create Conscious Machines. *arXiv:2403.00644* (2024).