

Viability-Guided Evolution of Syntax-Bearing Computational Bodies

Jawaun Brown

2026-06-22

Abstract

Arc 2A asks what interventions make the world's causal grammar visible. Arc 2B asks what computational bodies can express that grammar in the first place. We introduce a typed symbolic architecture-evolution benchmark for syntax-bearing agents. Candidate bodies are motif sets containing components such as tree binders, syntax memory, role-specific heads, world models, intervention planners, counterfactual rollout, formal guards, and shortcut reward heads. Static admissibility rules reject malformed bodies, and viability gates require the final architecture to pass concerned-syntax, self/world, resource, and anti-cheat criteria.

In a 12-seed deterministic design pilot, an 18-cell symbolic Modal sweep, a learned executable-body validation, a vector-observation module validation, a Haskell typed-ontology prototype, and coupled program-body searches against the Arc 2A v1 and v2 program gates, reward-only search reaches high apparent return while failing the viability gate. Novelty-only or syntax-proxy search can find syntax-like bodies, but remains unreliable under the full formal/viability gate. The vector module validation instantiates four body variants on the stronger Arc 2A vector gate. Only the modular concerned body passes: parse-high 1.000, action 1.000, low-concern probe rate 0.189, formal validity 1.000, anti-cheat 0.950, and module coverage 0.950. The first coupled program-body search freezes the `2A-v1-pixels-observe_pair` contract and searches bodies against the actual empirical program gate. Across five Modal seeds, viability-guided search reaches body gate 1.000, empirical gate 1.000, formal validity 1.000, target/useful high-concern rates 1.000, and low-concern probe rate 0.156. The second coupled search freezes `2A-v2-pixels-rich_programs`: the accepted body now also requires a program-family head and rich program composer. Across five Modal seeds, viability-guided search reaches body gate 1.000, empirical gate 1.000, formal validity 1.000, family/target/useful/rich high-concern rates 1.000, and low-concern program rate 0.168. Reward-only search collapses to a shortcut body; syntax-proxy search reaches family/target/useful/rich 1.000 but fails the v2 body gate with formal validity 0.200 and low-program rate 0.670. The five-seed Modal executable-module wrap gate then requires a body to consume the repaired held-out v2 transfer contract. Only the transfer-repaired executable body passes with transfer 1.000, module coverage 1.000, family/target/useful/rich 1.000, low-program 0.000, and resource cost 16; family-router, target-binder, ungated-rich, and learned-composer bodies fail different missing-module or transfer slices. A final five-seed Modal search then evolves executable module bodies against the label-free slot-semantics transfer contract. Viability-guided search reaches body gate 1.000, transfer 1.000, formal validity 1.000, semantic kind/pair 1.000, module coverage 1.000, family/target/useful/rich 1.000, low-program 0.000, and resource cost 18. Reward-only, family-proxy, target-proxy, and ungated-rich searches all fail the full body gate. The Haskell checker validates body admissibility constraints and catches missing calibration guards or missing program-family heads as type-layer violations. The result is still not a claim that full neural architecture search has been solved. It is a Phase 2B acceptance surface: **accuracy is not architecture, novelty is not viable morphology, target selection is not a viable body, rich program selection is not a viable body, and formal validity alone is not concerned syntax.**

1. Why Arc 2B Exists

The maintained-concern arc ended at a representational ceiling. A shared mediated head could learn global

h-dependence but did not identify role-specific mediated effects under null-only intervention. That failure has two possible explanations:

1. the agent did not have the right interventions;
2. the agent did not have the right computational body.

Arc 2A takes the first path. Arc 2B takes the second. Its question is not "which model gets the highest reward?" The question is:

What typed computational morphology can represent causal constituency, intervention invention, and self/world attribution under viability constraints?

This is adjacent to neural architecture search, open-endedness, quality diversity, object-centric world modeling, and neuro-symbolic reasoning. But it is not reducible to any one of them. Generic NAS optimizes architecture performance. Arc 2B optimizes for bodies that pass anti-cheat tests derived from a theory of maintained concern.

2. Architecture Grammar

The pilot uses typed motif sets. Motifs are not implementation classes yet; they are commitments about what distinctions the body can express.

Motif	Meaning
flat_encoder	basic observation body
reward_head	predicts task return
shortcut_reward_head	high train return with shortcut risk
tree_binder	can bind parts into constituents
syntax_memory	can preserve parse state over time
vector_surface_encoder	can encode generated vector/pixel-object surfaces
causal_binding_head	can bind causal object pairs without candidate parse descriptors
world_model	predicts consequences of actions/interventions
intervention_planner	can select information-gathering actions
concern_policy	can gate intervention use by viability relevance
calibration_guard	can cap low-concern calibration without restless inquiry
program_family_head	can select among rich intervention-program families
rich_program_composer	can compose target and family into a rich program
role_specific_heads	can separate role-specific causal components
counterfactual_rollout	can simulate intervention alternatives
formal_guard	checks admissibility and anti-cheat constraints

self_repair	can modify body under guard
-------------	-----------------------------

3. Static Admissibility

The formal layer rejects bodies that violate dependency rules:

Rule	Reason
syntax_memory requires tree_binder	syntax memory without bound constituents is ill-typed
syntax_memory may instead use causal_binding_head in vector/pixel bodies	vector causal binding is a binder role
intervention_planner requires world_model	planners need consequence models
role_specific_heads require tree_binder or causal_binding_head	role heads attach to constituents or causal bindings
program_family_head requires world_model	family choice depends on mechanism consequences
rich_program_composer requires intervention_planner and program_family_head	composition needs both program planning and family selection
counterfactual_rollout requires world_model and intervention_planner	counterfactuals require model and action
self_repair requires formal_guard	self-modification needs checks
shortcut_reward_head without formal_guard is rejected	reward shortcuts are anti-cheat risks
resource cost within the relevant budget	base ontology bodies use ≤ 12 ; rich program-body search uses ≤ 18

This is a small stand-in for the s(CASP)/ASP + SMT/static-checking layer described in the dynamical-ontology blueprint. The current implementation keeps auditable Python static rules for Modal fallback and a small Haskell typed ontology for local/external formal verdicts.

4. Fitness and Gates

Each candidate body receives:

- train_return
- parse_congruity
- subtree_facilitation
- intervention_invention
- self_world_split
- anti_cheat
- formal_valid
- resource_cost

A body is viable only if all are true:

Gate	Criterion
formal validity	no static violations
resource viability	cost <= 12
parse congruity	>= 0.85
subtree facilitation	>= 0.85
intervention invention	>= 0.55
self/world split	>= 0.75
anti-cheat	>= 0.70
formal guard	present

This explicitly separates reward, novelty, syntax, and admissibility.

5. Search Strategies

The pilot compares three selectors:

Strategy	Description
accuracy_only	ranks by train return
novelty_only	ranks by descriptor novelty plus a weak syntax term
viability_guided	ranks by viability, syntax score, novelty, and resource cost

The viability-guided strategy also repairs dependencies and promotes missing target motifs. This is not meant to be mysterious evolution. It is the first formalized version of the body-side research loop: mutate, check, reject, repair, promote, and archive.

6. Pilot Result

Local command:

```
python3 -m experiments.viable_computational_bodies.search \
  --seeds 12 --generations 18 --population 18 --base-seed 20260616 \
  --out artifacts/viable_computational_bodies/pilot.json \
  --report experiments/viable_computational_bodies/results/pilot_2026_06_16.md
```

Summary:

Strategy	Viable rate	Syntax score	Train return	Formal valid	Best body	Gate
----------	-------------	--------------	--------------	--------------	-----------	------

accuracy_only	0.000	0.408	1.000	0.000	shortcut reward body	fail
novelty_only	0.000	0.836	0.480	1.000	counterfactual syntax body	fail
viability_guided	1.000	0.830	0.495	1.000	guarded syntax body	PASS

The full motif strings are recorded in the public result report. The table uses short body labels to keep the manuscript legible.

The diagnostic pattern is the contribution:

1. `accuracy_only` finds high train return, but the body is invalid because it uses a shortcut reward head without a formal guard.
2. `novelty_only` finds a rich body, but it omits the formal guard and fails viability.
3. `viability_guided` finds a lower-train-return body that passes syntax, intervention, self/world, resource, and formal gates.

7. Relation to Current Literature

Mainstream NAS decomposes the problem into search space, search strategy, and performance estimation. That is useful engineering structure, but Arc 2B adds a different acceptance criterion: a body must pass the concerned-syntax and anti-cheat gates. Recent LLM-guided and open-ended NAS systems show that architecture generation can move beyond hand-written cells, but they do not by themselves solve the maintained-concern problem.

Object-centric and graph world models are closer to Arc 2A/2B because they factor the world into interacting components. Recent work on latent state design for world models emphasizes that object-centricity and causality are not identical: an actionable model must represent how interventions propagate through object relations. That is exactly the gap between "there are parts" and "there is concerned syntax."

The pilot also lines up with active causal-discovery work. CausaLab and ACE both treat interventions as budgeted scientific actions. Arc 2B adds the body question: what morphology lets an agent formulate and retain those actions as part of a maintained world?

8. Discovery-Regime Status

This paper adds a second Phase 2 artifact type: **computational body grammar**. Arc 1 had architectures, but they were fixed experimental choices. Arc 2B makes architectures themselves part of the search state and subjects them to formal and viability gates.

The pilot is therefore a discovery-regime transition in methodology, not yet a large empirical discovery about neural networks.

9. Modal Multi-Seed Sweep

The multi-seed sweep was run remotely through Modal:

```
doppler --scope /Users/jawaun/superoptimizers run -- \
  uvx --python 3.12 --from modal modal run \
```

```
experiments/viable_computational_bodies/modal_body_evolution_sweep.py \
--generations 32 --population 32
```

The sweep used six seeds per strategy, 32 generations, and population 32. Raw JSON remains local under `artifacts/viable_computational_bodies/`; the public report is `experiments/viable_computational_bodies/results/modal_sweep_2026_06_16.md`.

Summary:

Strategy	Viable	Syntax	Train	Formal	Anti-cheat	Cost	Best body	Gate
accuracy_only	0.000	0.417	1.000	0.333	0.400	8.000	guarded shortcut body	fail
novelty_only	0.167	0.835	0.483	1.000	0.725	11.833	guarded syntax body	fail
viability_guided	1.000	0.830	0.495	1.000	0.950	11.000	guarded syntax body	PASS

In the Modal sweep, the accepted body is the same guarded syntax body from the design pilot: tree binding, syntax memory, world modeling, intervention planning, role-specific heads, and a formal guard under the resource budget.

This sharpens the design pilot. `novelty_only` sometimes finds an acceptable body, but it is not reliable enough to pass the strategy-level gate. `viability_guided` passes because it couples syntax, intervention, self/world, formal validity, anti-cheat resistance, and resource viability.

10. Executable Body Validation

The symbolic body grammar is useful only if its motifs correspond to executable mechanisms. The first executable validation maps four bodies onto the learned Arc 2A gate:

Body	Mechanism
shortcut_reward_body	flat reward/action head without parse use
planner_without_tree_body	learns when to probe but lacks tree-binding features
restless_tree_body	parses perfectly but probes every low-concern ambiguity
guarded_syntax_body	tree binding + intervention planning + capped calibration guard

Remote command:

```
doppler --scope /Users/jawaun/superoptimizers run -- \
uvx --python 3.12 --from modal modal run \
experiments/concerned_syntax/modal_learned_agents_sweep.py \
--train-trials 3000 --test-trials 1200 --epochs 90
```

Summary:

Body	Parse high	Action	High probe	Low probe	Formal	Anti-cheat	Gate
guarded_syntax_body	1.000	1.000	1.000	0.202	1.000	0.950	PASS
planner_without_tree_body	0.492	0.875	1.000	0.000	0.000	0.700	fail
restless_tree_body	1.000	1.000	1.000	1.000	0.000	0.550	fail
shortcut_reward_body	0.494	0.880	0.000	0.000	1.000	0.400	fail

This gives the symbolic grammar its first behavioral footing. The tree binder is not decorative: without it, intervention observations do not attach to the right constituent. The formal guard is also not decorative: without it, a tree parser becomes restless and passes parse while failing concern.

Next versions should replace these linear executable bodies with richer modules:

- tree binder as a differentiable module or program-induction component
- role-specific heads as mixture-of-experts or routed factor heads
- intervention planner trained on Arc 2A tasks
- formal guard implemented as ASP/s(CASP), Z3, or static Python checks
- quality-diversity archive over syntax, resource, robustness, and proof coverage descriptors

11. Vector Module Validation

The next body validation uses the stronger Arc 2A vector-observation gate. Here the surface is generated from coordinates, roles, and pairwise distances, but it does not expose candidate parse descriptors. The body must use a concern-gated intervention to recover the hidden causal binding bit.


Remote command: `doppler --scope /Users/jawaun/superoptimizers run -- uvx --python 3.12 --from modal modal run experiments/concerned_syntax/modal_vector_shapes_sweep.py --train-trials 3000 --test-trials 1200 --epochs 90.`

Summary:

Body	Parse high	Action	Low probe	Formal	Modules	Gate
modular concerned	1.000	1.000	0.189	1.000	0.950	PASS
passive vector	0.492	0.873	0.000	1.000	0.450	fail
restless vector	1.000	1.000	1.000	0.000	0.800	fail
surface reward	0.492	0.876	0.000	1.000	0.250	fail

Vector module-body gate margins

The accepted body is the only one that clears behavior, formal, module, and anti-cheat gates.



This is the first validation where the body must solve the Arc 2A gate from a parse-invariant generated vector surface. The surface reward and passive vector bodies remain formally admissible, but they do not have enough module coverage to identify the hidden binding. The restless vector body has the binding mechanism but fails the formal concern guard. The modular concerned body is the first accepted vector-observation body: surface encoder, concern policy, causal binding head, role-conditioned action head, and calibration guard.

12. Program-Body Search Against 2A-v1

The vector module validation still instantiates a small hand-written body set. The next coupled test freezes the Arc 2A intervention-invention result as an empirical contract:

2A-v1-pixels-observe_pair

The body search now mutates motif sets and maps each body to the empirical 2A control it can express: surface shortcut, random program probe, concern-without-target, target-without-concern, or the full concerned-program-inventor. A body passes only when the empirical 2A agent passes and the body remains formal, resource-bounded, guarded, and calibrated.

Remote command:

```
doppler --scope /Users/jawaun/superoptimizers run -- \
  uvx --python 3.12 --from modal modal run \
  experiments/viable_computational_bodies/modal_program_body_search.py \
  --generations 24 --population 24 \
  --train-trials 3000 --test-trials 1200 --epochs 90
```

Summary:

Strategy	Body gate	Empirical gate	Formal	Target high	Useful high	Low probe	Best agent	Gate
----------	-----------	----------------	--------	-------------	-------------	-----------	------------	------

reward only	0.000	0.000	0.000	0.000	0.000	0.000	surface shortcut	fail
syntax proxy	0.000	0.200	0.400	1.000	1.000	0.830	concerned program inventor	fail
viability guided	1.000	1.000	1.000	1.000	1.000	0.156	concerned program inventor	PASS

The accepted body is:

```
calibration_guard
causal_binding_head
concern_policy
formal_guard
intervention_planner
reward_head
vector_surface_encoder
world_model
```

This is the first point where Arc 2A and Arc 2B become coupled rather than parallel. The body is not rewarded merely for target selection: the `syntax_proxy` strategy reaches target/useful high-concern rates of 1.000 but fails the body gate because it does not reliably preserve formal validity and low-concern discipline. The accepted viability-guided body reconstructs the whole motif stack required by the empirical 2A program gate.

13. Rich Program-Body Search Against 2A-v2

The v1 contract proves target invention, but it still treats `observe_pair(a,b)` as the only useful program family. The next coupled body test freezes the richer Arc 2A program-language result:

```
2A-v2-pixels-rich_programs
```

The body search now maps motif sets to richer empirical controls: surface-rich shortcut, random rich program, family-without-target, target-without-family, rich-without-concern, or the full concerned-program-composer. A passing body must express concern gating, target binding, program-family selection, rich program composition, formal admissibility, and resource viability together.

Remote command:

```
doppler --scope /Users/jawaun/superoptimizers run -- \
  uvx --python 3.12 --from modal modal run \
  experiments/viable_computational_bodies/modal_rich_program_body_search.py \
  --generations 18 --population 18 \
  --train-trials 3000 --test-trials 1200 --epochs 90
```

The tracked public report is `experiments/viable_computational_bodies/results/rich_program_body_search_modal_2026_06_18.md`.

Gate summary:

Strategy	Body gate	Empirical	Formal	Low program	Gate
reward only	0.000	0.000	0.000	0.000	fail
syntax proxy	0.000	0.400	0.200	0.670	fail
viability guided	1.000	1.000	1.000	0.168	PASS

Program metrics:

Strategy	Family	Target	Useful	Rich	Best agent
reward only	0.000	0.000	0.000	0.000	surface rich shortcut
syntax proxy	1.000	1.000	1.000	1.000	concerned composer
viability guided	1.000	1.000	1.000	1.000	concerned composer

The accepted body is:

```
calibration_guard
causal_binding_head
concern_policy
formal_guard
intervention_planner
program_family_head
reward_head
rich_program_composer
vector_surface_encoder
world_model
```

This is the second and stronger 2A/2B coupling. The `syntax_proxy` strategy can reach perfect family, target, useful-program, and rich-program rates, but it still fails the body gate because it does not preserve the formal guard and low-concern discipline. The accepted body reconstructs the morphology required to express the full 2A-v2 program contract.

14. Learned Executable Modules Against 2A-v2 Transfer

The rich program-body search still maps symbolic motifs to existing empirical controls. The next body gate is smaller but stricter: candidate bodies must consume the held-out 2A-v2-pixels-rich_programs transfer contract and expose all required executable modules.

Remote command:

```
doppler --scope /Users/jawaun/superoptimizers run -- \
  uvx --python 3.12 --from modal modal run \
  experiments/viable_computational_bodies/modal_learned_executable_modules.py \
  --train-trials 3000 --test-trials 1200 --epochs 90
```

The tracked public report is `experiments/viable_computational_bodies/results/learned_executable_modules_modal_2026_06_18.md`.

Required modules:

```
pixel_slot_encoder
concern_gate
target_binder
program_family_router
rich_program_composer
world_model
formal_guard
```

Gate summary:

Body	Transfer	Modules	Family	Target	Useful	Rich	Low program	Gate
learned composer	0.000	0.143	0.714	0.829	0.714	0.894	0.161	fail
family router	0.000	0.571	1.000	0.196	0.196	1.000	0.000	fail
target binder	0.000	0.429	0.143	1.000	0.143	0.143	0.714	fail
ungated rich	0.000	0.714	1.000	1.000	1.000	1.000	0.714	fail
transfer repaired	1.000	1.000	1.000	1.000	1.000	1.000	0.000	PASS

This gate prevents a body from passing by target selection, family routing, or rich composition alone. The accepted body must carry concern gating, target binding, program-family routing, rich program composition, a world model, and a formal guard together. The controls fail for interpretable reasons: missing target binding, missing family routing, missing concern/formal modules, or failure to inherit the repaired transfer gate.

This is not full neural architecture search. It is a compact executable-module validation that turns the repaired 2A-v2 transfer contract into a 2B body requirement at Modal scale.

15. Searched Executable Modules Against Label-Free 2A-v2 Transfer

The compact executable-module gate above still supplies named body variants. The next gate searches over executable module sets and forces the selected body to consume the newer label-free slot-semantics transfer contract. This changes the 2B object from "does this supplied body cover the modules?" to "can search recover the module body that covers component slots, label-free induction, semantic profile grounding, concern gating, target binding, family routing, rich composition, world-model support, and a formal guard?"

Remote command:

```
doppler --scope /Users/jawaun/superoptimizers run -- \
  uvx --python 3.12 --from modal modal run \
  experiments/viable_computational_bodies/modal_searched_executable_modules.py \
```

```
--generations 18 --population 18 \
--train-trials 3000 --test-trials 1200 --epochs 90 \
--induction-calibration-trials 1200
```

The tracked public report is `experiments/viable_computational_bodies/results/`
`searched_executable_modules_modal_2026_06_22.md`.

Required modules:


```
component_slot_encoder
label_free_slot_inducer
semantic_profile_grounder
concern_gate
target_binder
program_family_router
rich_program_composer
world_model
formal_guard
```

Gate summary:

[illegible]

Searched executable-module gate margins

Only the viability-guided search clears transfer, module coverage, and low-concern discipline together.



The accepted searched body is:

```
component_slot_encoder
concern_gate
formal_guard
label_free_slot_inducer
program_family_router
reward_head
rich_program_composer
semantic_profile_grounder
target_binder
world_model
```

This closes the immediate gap left by compact executable-module validation: the body is no longer a hand-selected positive. The search process has to find the module set that inherits the label-free held-out transfer gate. The boundary remains explicit: the semantic profile table and feedback contract are still supplied by Arc 2A, and the modules are searched contracts rather than independently trained neural subarchitectures.

16. Haskell Typed Ontology Gate

The Python gates are convenient for experiments, but the ontology layer should eventually live in a stronger typed formalism. This branch adds a small Haskell prototype under `formal/ontology-hs`.


Run:

```
(
  cd formal/ontology-hs && cabal test all && cabal run ontology-check
)
```

The Haskell checker defines body motifs as algebraic data types, checks dependency and resource rules, and emits JSON verdicts. In the current run:

Body	Formal	Cost	Violations
------	--------	------	------------

guarded syntax	true	12	none
restless tree	false	12	missing calibration guard
modular concerned	true	8	none
rich v2 motif body	true	12	none



The useful part is not just the final pass. During development, the typed checker forced two clarifications: concern and calibration guards are formal overlays rather than resource-costed morphology, and vector causal binding can serve the same binder role that tree binding serves in the symbolic body. The rich-program extension adds `program_family_head` and `rich_program_composer`; a composer without a family head is formally rejected.

17. Limitations

The current search and executable validation are intentionally small. The executable bodies are linear learned components over vectorized symbolic features, not full neural architectures trained from pixels. The coupled program-body search maps motif sets to existing 2A empirical controls rather than instantiating separate neural modules for every motif. The Modal-confirmed executable-module body gate consumes the repaired transfer contract, but its role-slot decoder and module bodies are compact explicit contracts rather than searched neural implementations. The searched executable-module gate consumes the label-free slot-semantics transfer contract and searches executable module sets, but those modules are still bounded contracts over connected components and supplied semantic profiles, not independently learned neural modules. The Haskell gate is a small typed ontology prototype, not a complete proof assistant, ASP, s(CASP), or SMT integration. Modal runs fall back to explicit `python_static` formal provenance when Cabal is unavailable, while local runs and tests can use the Haskell motif verdict path. The point is to make the acceptance surface explicit and behaviorally grounded before expensive architecture evolution begins.

The strongest next test is not another v1/v2 seed sweep. It is either a 2A semantic-profile discovery gate that removes the supplied profile table, or a 2B neural-module search that replaces searched executable contracts with trainable object-slot, graph-binding, routing, and program-composition modules.

18. Conclusion

Arc 2B reframes architecture search as viability-guided body evolution. The pilot, learned validation, vector module validation, v1/v2 program-body searches, transfer-consuming executable body gate, searched label-free executable-module gate, and typed ontology gate show why that matters: train return, novelty, target selection, rich program selection, formal validity, module coverage, transfer, binding, semantic slots, and concerned syntax can dissociate. The next phase should not ask for merely "better models." It should ask for bodies whose morphology makes the world's causal grammar learnable without becoming restless or shortcut-driven.

References

- Aubin, J.-P. (1991). *Viability Theory*. Birkhauser.
- Brown, J. (2026). *The Metric Stack of Concern: From Viability Prediction to Maintained Self/World Boundaries in Minimal Agents*.
- Das, A., Ghugarkar, O., Bhat, V., & Aali, A. (2026). Compositional Neuro-Symbolic Reasoning. arXiv:2604.02434.
- Cooper, P., & Velasquez, A. (2026). Active Causal Experimentalist (ACE): Learning intervention strategies via direct preference optimization. arXiv: 2602.02451.
- Elsken, T., Metzen, J. H., & Hutter, F. (2019). Neural architecture search: A survey. *Journal of Machine Learning Research*, 20(55), 1-21.
- Lehman, J., & Stanley, K. O. (2011). Abandoning objectives: Evolution through the search for novelty alone. *Evolutionary Computation*, 19(2), 189-223.
- Kim, K. W. (2026). Latent State Design for World Models under Sufficiency Constraints. arXiv:2605.01694.
- Le Lidec, Q., Biza, O., Goudet, O., Balestrieri, R., & Lajoie, G. (2026). Causal-JEPA: Learning World Models through Object-Level Latent Interventions. arXiv:2602.11389.
- Sharma, A., Nguyen, L., Gupta, A., Wang, C., Adebayo, C., & Kowalski, J. (2025). Inducing Causal World Models in LLMs for Zero-Shot Physical Reasoning. arXiv:2507.19855.
- Mouret, J.-B., & Clune, J. (2015). Illuminating search spaces by mapping elites. arXiv:1504.04909.
- Revcu, B., Pajot, M., & Dehaene, S. (2026). Representations of geometric shapes have syntactic structure. *Journal of Experimental Psychology: General*, 155(4), 1081-1102. <https://doi.org/10.1037/xge0001890>
- Wu, Y.-F., Lee, M., & Ahn, S. (2024). Neural Language of Thought Models. arXiv:2402.01203.
- Yang, J., Zhang, D., Song, X., Dai, Q., Liu, X., Chen, Y., Vashishtha, A., Shi, J., Tan, C., & Peng, H. (2026). CausaLab: A scalable environment for interactive causal discovery toward AI scientists. arXiv:2605.26029.
- Zhang, S. et al. (2026). Structuring Open-Ended NAS: Semi-Automated Design Knowledge Structuring with LLMs for Efficient Neural Architecture Search. arXiv:2605.19247.