

Viability-Guided Evolution of Syntax-Bearing Computational Bodies

Jawaun Brown

2026-06-16

Abstract

Arc 2A asks what interventions make the world's causal grammar visible. Arc 2B asks what computational bodies can express that grammar in the first place. We introduce a typed symbolic architecture-evolution benchmark for syntax-bearing agents. Candidate bodies are motif sets containing components such as tree binders, syntax memory, role-specific heads, world models, intervention planners, counterfactual rollout, formal guards, and shortcut reward heads. Static admissibility rules reject malformed bodies, and viability gates require the final architecture to pass concerned-syntax, self/world, resource, and anti-cheat criteria.

In a 12-seed deterministic design pilot, an 18-cell symbolic Modal sweep, a learned executable-body validation, a vector-observation module validation, and a Haskell typed-ontology prototype, reward-only search reaches high apparent return while failing the viability gate. Novelty-only search finds syntax-like bodies, but remains unreliable under the full formal/viability gate. The vector module validation instantiates four body variants on the stronger Arc 2A vector gate. Only the modular concerned body passes: parse-high 1.000, action 1.000, low-concern probe rate 0.189, formal validity 1.000, anti-cheat 0.950, and module coverage 0.950. The Haskell checker separately validates body admissibility constraints and catches missing calibration guards as type-layer violations. The result is still not a claim that full neural architecture search has been solved. It is a Phase 2B acceptance surface: **accuracy is not architecture, novelty is not viable morphology, and formal validity alone is not concerned syntax.**

1. Why Arc 2B Exists

The maintained-concern arc ended at a representational ceiling. A shared mediated head could learn global h-dependence but did not identify role-specific mediated effects under null-only intervention. That failure has two possible explanations:

1. the agent did not have the right interventions;
2. the agent did not have the right computational body.

Arc 2A takes the first path. Arc 2B takes the second. Its question is not "which model gets the highest reward?" The question is:

What typed computational morphology can represent causal constituency, intervention invention, and self/world attribution under viability constraints?

This is adjacent to neural architecture search, open-endedness, quality diversity, object-centric world modeling, and neuro-symbolic reasoning. But it is not reducible to any one of them. Generic NAS optimizes architecture performance. Arc 2B optimizes for bodies that pass anti-cheat tests derived from a theory of maintained concern.

2. Architecture Grammar

The pilot uses typed motif sets. Motifs are not implementation classes yet; they are commitments about what distinctions the body can express.

Motif	Meaning
flat_encoder	basic observation body
reward_head	predicts task return
shortcut_reward_head	high train return with shortcut risk
tree_binder	can bind parts into constituents
syntax_memory	can preserve parse state over time
world_model	predicts consequences of actions/interventions
intervention_planner	can select information-gathering actions
role_specific_heads	can separate role-specific causal components
counterfactual_rollout	can simulate intervention alternatives
formal_guard	checks admissibility and anti-cheat constraints
self_repair	can modify body under guard

3. Static Admissibility

The formal layer rejects bodies that violate dependency rules:

Rule	Reason
syntax_memory requires tree_binder	syntax memory without bound constituents is ill-typed
intervention_planner requires world_model	planners need consequence models
role_specific_heads require tree_binder	role heads attach to constituents
counterfactual_rollout requires world_model and intervention_planner	counterfactuals require model and action
self_repair requires formal_guard	self-modification needs checks
shortcut_reward_head without formal_guard is rejected	reward shortcuts are anti-cheat risks
resource cost ≤ 12	bodies must remain viable

This is a small stand-in for the s(CASP)/ASP + SMT/static-checking layer described in the dynamical-ontology blueprint. The current implementation is plain Python so the gate is auditable and easy to test.

4. Fitness and Gates

Each candidate body receives:

- `train_return`

- parse_congruity
- subtree_facilitation
- intervention_invention
- self_world_split
- anti_cheat
- formal_valid
- resource_cost

A body is viable only if all are true:

Gate	Criterion
formal validity	no static violations
resource viability	cost <= 12
parse congruity	>= 0.85
subtree facilitation	>= 0.85
intervention invention	>= 0.55
self/world split	>= 0.75
anti-cheat	>= 0.70
formal guard	present

This explicitly separates reward, novelty, syntax, and admissibility.

5. Search Strategies

The pilot compares three selectors:

Strategy	Description
accuracy_only	ranks by train return
novelty_only	ranks by descriptor novelty plus a weak syntax term
viability_guided	ranks by viability, syntax score, novelty, and resource cost

The viability-guided strategy also repairs dependencies and promotes missing target motifs. This is not meant to be mysterious evolution. It is the first formalized version of the body-side research loop: mutate, check, reject, repair, promote, and archive.

6. Pilot Result

Local command:

```
python3 -m experiments.viable_computational_bodies.search \
  --seeds 12 --generations 18 --population 18 --base-seed 20260616 \
  --out artifacts/viable_computational_bodies/pilot.json \
  --report experiments/viable_computational_bodies/results/pilot_2026_06_16.md
```

Summary:

Strategy	Viable rate	Syntax score	Train return	Formal valid	Best body	Gate
accuracy_only	0.000	0.408	1.000	0.000	shortcut reward body	fail
novelty_only	0.000	0.836	0.480	1.000	counterfactual syntax body	fail
viability_guided	1.000	0.830	0.495	1.000	guarded syntax body	PASS

The full motif strings are recorded in the public result report. The table uses short body labels to keep the manuscript legible.

The diagnostic pattern is the contribution:

1. `accuracy_only` finds high train return, but the body is invalid because it uses a shortcut reward head without a formal guard.
2. `novelty_only` finds a rich body, but it omits the formal guard and fails viability.
3. `viability_guided` finds a lower-train-return body that passes syntax, intervention, self/world, resource, and formal gates.

7. Relation to Current Literature

Mainstream NAS decomposes the problem into search space, search strategy, and performance estimation. That is useful engineering structure, but Arc 2B adds a different acceptance criterion: a body must pass the concerned-syntax and anti-cheat gates. Recent LLM-guided and open-ended NAS systems show that architecture generation can move beyond hand-written cells, but they do not by themselves solve the maintained-concern problem.

Object-centric and graph world models are closer to Arc 2A/2B because they factor the world into interacting components. Recent work on latent state design for world models emphasizes that object-centricity and causality are not identical: an actionable model must represent how interventions propagate through object relations. That is exactly the gap between "there are parts" and "there is concerned syntax."

The pilot also lines up with active causal-discovery work. CausaLab and ACE both treat interventions as budgeted scientific actions. Arc 2B adds the body question: what morphology lets an agent formulate and retain those actions as part of a maintained world?

8. Discovery-Regime Status

This paper adds a second Phase 2 artifact type: **computational body grammar**. Arc 1 had architectures, but they were fixed experimental choices. Arc 2B makes architectures themselves part of the search state and subjects them to formal and viability gates.

The pilot is therefore a discovery-regime transition in methodology, not yet a large empirical discovery about neural networks.

9. Modal Multi-Seed Sweep

The multi-seed sweep was run remotely through Modal:

```
doppler --scope /Users/jawaun/superoptimizers run -- \
  uvx --python 3.12 --from modal modal run \
  experiments/viable_computational_bodies/modal_body_evolution_sweep.py \
  --generations 32 --population 32
```

The sweep used six seeds per strategy, 32 generations, and population 32. Raw JSON remains local under `artifacts/viable_computational_bodies/`; the public report is `experiments/viable_computational_bodies/results/modal_sweep_2026_06_16.md`.

Summary:

Strategy	Viable	Syntax	Train	Formal	Anti-cheat	Cost	Best body	Gate
accuracy_only	0.000	0.417	1.000	0.333	0.400	8.000	guarded shortcut body	fail
novelty_only	0.167	0.835	0.483	1.000	0.725	11.833	guarded syntax body	fail
viability_guided	1.000	0.830	0.495	1.000	0.950	11.000	guarded syntax body	PASS

In the Modal sweep, the accepted body is the same guarded syntax body from the design pilot: tree binding, syntax memory, world modeling, intervention planning, role-specific heads, and a formal guard under the resource budget.

This sharpens the design pilot. `novelty_only` sometimes finds an acceptable body, but it is not reliable enough to pass the strategy-level gate. `viability_guided` passes because it couples syntax, intervention, self/world, formal validity, anti-cheat resistance, and resource viability.

10. Executable Body Validation

The symbolic body grammar is useful only if its motifs correspond to executable mechanisms. The first executable validation maps four bodies onto the learned Arc 2A gate:

Body	Mechanism
shortcut_reward_body	flat reward/action head without parse use
planner_without_tree_body	learns when to probe but lacks tree-binding features
restless_tree_body	parses perfectly but probes every low-concern ambiguity
guarded_syntax_body	tree binding + intervention planning + capped calibration guard

Remote command:

```
doppler --scope /Users/jawaun/superoptimizers run -- \
  uvx --python 3.12 --from modal modal run \
  experiments/concerned_syntax/modal_learned_agents_sweep.py \
  --train-trials 3000 --test-trials 1200 --epochs 90
```

Summary:

Body	Parse high	Action	High probe	Low probe	Formal	Anti-cheat	Gate
guarded_syntax_body	1.000	1.000	1.000	0.202	1.000	0.950	PASS
planner_without_tree_body	0.492	0.875	1.000	0.000	0.000	0.700	fail
restless_tree_body	1.000	1.000	1.000	1.000	0.000	0.550	fail
shortcut_reward_body	0.494	0.880	0.000	0.000	1.000	0.400	fail

This gives the symbolic grammar its first behavioral footing. The tree binder is not decorative: without it, intervention observations do not attach to the right constituent. The formal guard is also not decorative: without it, a tree parser becomes restless and passes parse while failing concern.

Next versions should replace these linear executable bodies with richer modules:

- tree binder as a differentiable module or program-induction component
- role-specific heads as mixture-of-experts or routed factor heads
- intervention planner trained on Arc 2A tasks
- formal guard implemented as ASP/s(CASP), Z3, or static Python checks
- quality-diversity archive over syntax, resource, robustness, and proof coverage descriptors

11. Vector Module Validation

The next body validation uses the stronger Arc 2A vector-observation gate. Here the surface is generated from coordinates, roles, and pairwise distances, but it does not expose candidate parse descriptors. The body must use a concern-gated intervention to recover the hidden causal binding bit.


Remote command: `doppler --scope /Users/jawaun/superoptimizers run -- uvx --python 3.12 --from modal modal run experiments/concerned_syntax/modal_vector_shapes_sweep.py --train-trials 3000 --test-trials 1200 --epochs 90.`

Summary:

Body	Parse high	Action	Low probe	Formal	Modules	Gate
modular concerned	1.000	1.000	0.189	1.000	0.950	PASS
passive vector	0.492	0.873	0.000	1.000	0.450	fail
restless vector	1.000	1.000	1.000	0.000	0.800	fail
surface reward	0.492	0.876	0.000	1.000	0.250	fail

Vector module-body gate margins

The accepted body is the only one that clears behavior, formal, module, and anti-cheat gates.



This is the first validation where the body must solve the Arc 2A gate from a parse-invariant generated vector surface. The surface reward and passive vector bodies remain formally admissible, but they do not have enough module coverage to identify the hidden binding. The restless vector body has the binding mechanism but fails the formal concern guard. The modular concerned body is the first accepted vector-observation body: surface encoder, concern policy, causal binding head, role-conditioned action head, and calibration guard.

12. Haskell Typed Ontology Gate


The Python gates are convenient for experiments, but the ontology layer should eventually live in a stronger typed formalism. This branch adds a small Haskell prototype under `formal/ontology-hs`.

Run:

```
(
  cd formal/ontology-hs && cabal test all && cabal run ontology-check
)
```

The Haskell checker defines body motifs as algebraic data types, checks dependency and resource rules, and emits JSON verdicts. In the current run:

Body	Formal	Cost	Violations
guarded syntax	true	12	none
restless tree	false	12	missing calibration guard
modular concerned	true	8	none



The useful part is not just the final pass. During development, the typed checker forced two clarifications: concern and calibration guards are formal overlays rather than resource-costed morphology, and vector causal binding can serve the same binder role that tree binding serves in the symbolic body.

13. Limitations

The current search and executable validation are intentionally small. The executable bodies are linear learned components over vectorized symbolic features, not full neural architectures trained from pixels. The Haskell gate is a small typed ontology prototype, not a complete proof assistant, ASP, s(CASP), or SMT integration. The point is to make the acceptance surface explicit and behaviorally grounded before expensive architecture evolution begins.

The strongest next test is whether executable bodies discovered under this grammar pass the Arc 2A pixel-level concerned-syntax benchmark without direct access to the hidden parse.

14. Conclusion

Arc 2B reframes architecture search as viability-guided body evolution. The pilot, learned validation, vector module validation, and typed ontology gate show why that matters: train return, novelty, formal validity, module coverage, binding, and concerned syntax can dissociate. The next phase should not ask for merely "better models." It should ask for bodies whose morphology makes the world's causal grammar learnable without becoming restless or shortcut-driven.

References

- Aubin, J.-P. (1991). *Viability Theory*. Birkhauser.
- Brown, J. (2026). *The Metric Stack of Concern: From Viability Prediction to Maintained Self/World Boundaries in Minimal Agents*.
- Das, A., Ghugarkar, O., Bhat, V., & Aali, A. (2026). Compositional Neuro-Symbolic Reasoning. arXiv:2604.02434.
- Cooper, P., & Velasquez, A. (2026). Active Causal Experimentalist (ACE): Learning intervention strategies via direct preference optimization. arXiv: 2602.02451.

- Elsken, T., Metzen, J. H., & Hutter, F. (2019). Neural architecture search: A survey. *Journal of Machine Learning Research*, 20(55), 1-21.
- Lehman, J., & Stanley, K. O. (2011). Abandoning objectives: Evolution through the search for novelty alone. *Evolutionary Computation*, 19(2), 189-223.
- Kim, K. W. (2026). Latent State Design for World Models under Sufficiency Constraints. arXiv:2605.01694.
- Le Lidec, Q., Biza, O., Goudet, O., Balestriero, R., & Lajoie, G. (2026). Causal-JEPA: Learning World Models through Object-Level Latent Interventions. arXiv:2602.11389.
- Sharma, A., Nguyen, L., Gupta, A., Wang, C., Adebayo, C., & Kowalski, J. (2025). Inducing Causal World Models in LLMs for Zero-Shot Physical Reasoning. arXiv:2507.19855.
- Mouret, J.-B., & Clune, J. (2015). Illuminating search spaces by mapping elites. arXiv:1504.04909.
- Revenu, B., Pajot, M., & Dehaene, S. (2026). Representations of geometric shapes have syntactic structure. *Journal of Experimental Psychology: General*, 155(4), 1081-1102. <https://doi.org/10.1037/xge0001890>
- Wu, Y.-F., Lee, M., & Ahn, S. (2024). Neural Language of Thought Models. arXiv:2402.01203.
- Yang, J., Zhang, D., Song, X., Dai, Q., Liu, X., Chen, Y., Vashishtha, A., Shi, J., Tan, C., & Peng, H. (2026). CausaLab: A scalable environment for interactive causal discovery toward AI scientists. arXiv:2605.26029.
- Zhang, S. et al. (2026). Structuring Open-Ended NAS: Semi-Automated Design Knowledge Structuring with LLMs for Efficient Neural Architecture Search. arXiv:2605.19247.